

1


This slide has a light blue background with a subtle wave pattern. In the top left corner, there is an orange icon of a database cylinder. The title "ERM – Origin and Role" is centered in a bold, orange font. Below the title, there are two bullet points in a dark red font. The first bullet point discusses the introduction of the ERM by Chen and its formalization by Ng, with a sub-point mentioning Chen's original paper. The second bullet point discusses the use of ERM in structured methodologies, with a sub-point giving examples like LDS, ELH, and Dataflow. At the bottom of the slide, there is a footer with three items: "Joseph Vella - Introduction to Databases" on the left, the number "2" in the center, and "Logical Database Design" on the right.

ERM – Origin and Role

- One of the most popular high level conceptual data modelling language is the *Entity Relationship Model (ERM)* introduced by **Chen** and formalised by **Ng**.
 - For the record Chen's original paper is still considered as "one of the most influential papers in the computer software field".
- In many *structured methodologies* for application development the ERM is used in complementary fashion with "functional" descriptions and "data flow" paths.
 - E.g. LDS, ELH & Dataflow

Joseph Vella - Introduction to Databases 2 Logical Database Design

2




Entity-Relationship Modelling

- With the Entity-Relationship Modelling techniques, a database designer can reshape the application domain at a conceptual (high) level.
 - The main aim of the ERM is the enumeration (e.g., listing) of the application domain's entities and their inter relationships.
 - The detail level of an ER model depends on the scope of the application.
 - The ERM representation is graphical (e.g., graph, model, or diagram).
 - The ERM is "value based".
 - ERM are independent of any database management system.

Joseph Vella - Introduction to Databases 3 Logical Database Design

3

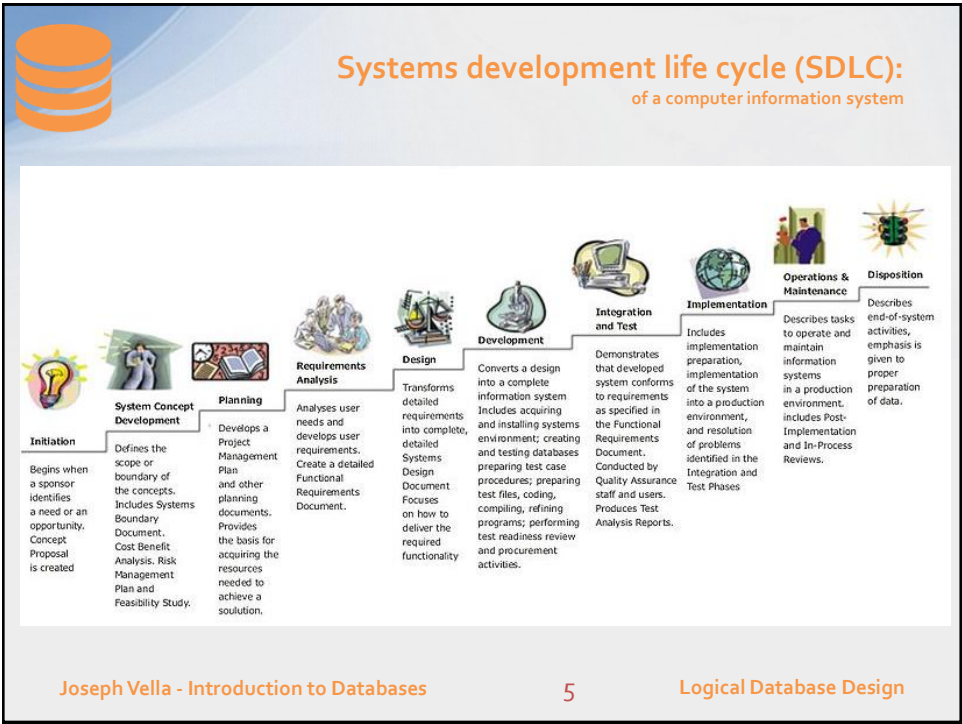


High Level Database Design


- Clearly identify the application domain:
 - First cut listing of users' requirements;
 - Analysis of domain and its requirements.
- Result:
 - Isolation of the information system entities, and their relationships.
 - The main functions of the system are supplied by these entities and relationship instances; and
 - A list of users' requirements (forms upkeep, queries, reports generation, etc.).

Joseph Vella - Introduction to Databases 4 Logical Database Design

4



5



SDLC and DBLC overlap

	Initiation	System Concept Development	Planning	Requirements Analysis	Design	Development	Integration and Test	Implementation	Operations and Maintenance	Disposition
Information Requirements				+					+	
Logical Design					+				+	+
Physical Design					+	+	+		+	+
Implementation						+	+	+	+	

Joseph Vella - Introduction to Databases

6

Logical Database Design

6



Systems Development Risks


- Better not to underestimate the initial stages (e.g., requirements and data requirements);
- Too formal against too informal specification dilemma;
- Technological advancement not properly gauged during design and development;
- Changing requirements. For example, in Legal; Business environment.

Joseph Vella - Introduction to Databases

7









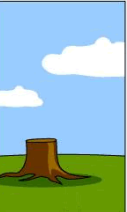

Logical Database Design

7



Systems Development Risks

a case study -cartoon


 <p>How the customer explained it</p>	 <p>How the Project Leader understood it</p>	 <p>How the Analyst designed it</p>	 <p>How the Programmer wrote it</p>	 <p>How the Business Consultant described it</p>
 <p>How the project was documented</p>	 <p>What operations installed</p>	 <p>How the customer was billed</p>	 <p>How it was supported</p>	 <p>What the customer really needed</p>

Joseph Vella - Introduction to Databases

8

Logical Database Design

8




Practical Uses of Conceptual Design

- **Starting point** for the *database design phase* (the DBMS design is **package dependent**).
 - Since the ERM language is *data model independent*, therefore different procedures are required to map the high-level conceptual design onto a particular data model's constructs.
- **Verifying** that the users' requirements are met.
- **Validating** that the users' requirements are consistent.
- **Test base** for the application's functional requirements and query design (more specific than those present in the initial application domain specification).
- **View integration** for multi-databases and front-end development.
- **Note:**
If a deficiency is found, for example a query cannot be answered from the present ERM, an alteration to the design is mandatory!

Joseph Vella - Introduction to Databases 9 Logical Database Design


9



Entities

Joseph Vella - Introduction to Databases 10 Logical Database Design

10




Entities

- An **entity** is a named construct representing a real or abstract player in the application's domain of discourse.
It is expected that an entity's instances have an **independent existence** in the discourse.
 - E.g. an EMPLOYEE is an entity in a projects database.
- An entity is depicted as a **named rectangle** in the model's graph.
- ✓ Each **entity name** is unique.
- ✓ It is expected that each entity has many **instances**.
 - ✓ Note: sometimes we represent these as $entity_n$. But instances are not drawn on the diagram!

EMPLOYEE

Joseph Vella - Introduction to Databases
11
Logical Database Design

11



Entity Attributes

- An entity is also described by a number of properties called **attributes**.
 - An entity is said to be **composed** of its attributes.
- Each attribute has a **name** and is annotated (but not shown) by a **value domain**.
 - E.g. in EMPLOYEE entity typical attributes are *Name*, *Address*, *Date_of_Birth*, and *Salary*.
 - With value domains of *string*, *string*, *date* and *real number* respectively.
- Each attribute is depicted as a **named oval** (or a rectangle with softened edges) attached to its emanating entity.

EMPLOYEE

Name

Address


Date Of Birth

Salary

- ✓ Attribute names have to be unique within an entity.
- ✓ Entity instances would then have to be assigned values to these attributes from the associated value domains.
 - ✓ i.e. ERM instances are **value based**.

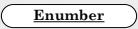
Joseph Vella - Introduction to Databases
12
Logical Database Design

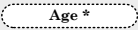
12



Entity Attribute Flavours (i)


- **Primary key set**
 - Is an attribute (or set of attributes) whose value(s) uniquely distinguish an instance from all others from the same entity.
 - **There can be only one primary key set depiction per entity.**
 - We depict this constraint by *underlining* the attribute names in the key set.
- **Computed Attribute (vs Valued Attribute)**
 - Some attributes values is directly derived from other values present in the same instance - **computed**.
 - **E.g. given *Date_of_Birth* value then it's easy to compute his current *Age*.**
 - We depict this type of attribute by drawing a *named dotted oval*.






Joseph Vella - Introduction to Databases 13 Logical Database Design

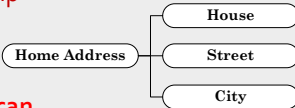
13



Entity Attribute Flavours (ii)


- **Single vs Multi-valued attributes**
 - Another aspect of attributes is the allowed number of values to be assigned to an entity's property.
 - In some other cases an attribute takes many values from its associated domain - a **multi-valued** attribute and these are drawn with *double ovals*.
 - **For example, employee's *Qualifications* is a multi-valued attribute.**
- **Simple vs. Composite**
 - Sometimes a sub-set of an entity's attributes tend to have a high affinity and it's advantageous to lump these together into what is called a **composite** attribute.
 - **For example, assume we have the following attributes, *House*, *Street* and *City*. These three can form the composite attribute *Home_address*.**



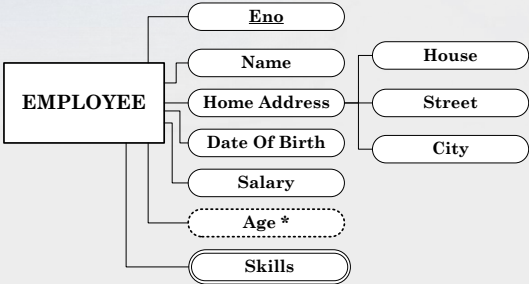


Joseph Vella - Introduction to Databases 14 Logical Database Design

14



Entity Example



```
graph LR; EMPLOYEE[EMPLOYEE] --- Eno[Eno]; EMPLOYEE --- Name[Name]; EMPLOYEE --- HomeAddress[Home Address]; EMPLOYEE --- DOB[Date Of Birth]; EMPLOYEE --- Salary[Salary]; EMPLOYEE --- Age["Age *"]; EMPLOYEE --- Skills[Skills]; HomeAddress --- House[House]; HomeAddress --- Street[Street]; HomeAddress --- City[City];
```

*) Age is current date less DateOfBirth.

Joseph Vella - Introduction to Databases

15

Logical Database Design

15




Relationships

Joseph Vella - Introduction to Databases

16


Logical Database Design

16




Relationships

- A **relationship** is a named association between a number of entities that is governed by certain rules and sometimes qualified with its own attributes.
 - A **relationship instance** is a particular case of relationship and qualified by the respective entity instances it's relating.
 - A relationship name is unique between its participating entities.
 - A relationship is depicted with:
 - a **diamond that inscribes its name**;
 - a **line connecting the relationship (diamond) with each participating entity (rectangle)**. This line carries two modeling indications:
 - **Cardinality** (number – e.g., 1 and N);
 - **Participation** (single or double line).



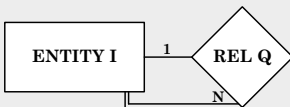
Joseph Vella - Introduction to Databases 17 Logical Database Design

17




Relationships – Degree constraints (i)

- This constraint deals with enumerating the entities in a relationship.
 - Any relationship instance must have an entity instance from each respective participating entities.
- The **degree** of a relationship is the number of entities that define the relationship.
 - If there are two, then the relationship is **binary** and if there are three then it's called **ternary** etc.
 - If more than three we call them **n-ary**.
 - In practice the most common type is the binary relationship.
- If an entity is repeated in a relationship definition then the relationship is said to be **recursive**.






Joseph Vella - Introduction to Databases 18 Logical Database Design

18



Relationships – Cardinality constraints (ii)


- How many times can an entity's instances appear in separate in relationship instances?
 - The most common cardinalities are the "one" (for once) to "many".
- For a binary relationship the number of these possibilities are four namely;
 - one to one (1-1),
 - one to many (1-N),
 - many to one (N-1), and
 - many to many (N-M).

- In the modelling diagram the cardinalities are depicted by including the symbols "1" or "N" (or "M", "P") on the opposite edges of the relationship diamond.


Joseph Vella - Introduction to Databases 19 Logical Database Design

19



Relationships – Participation constraints (iii)


- Another constraint is the level of participation of an entity's instances in a relationship.
- The options are **total** (or **mandatory**) and **partial** (or **optional**) for each of the participating entity.
 - Total participation implies that all instances must participate in at least one relationship instance.
 - Partial participation implies that an entity instance may not be present in any relationship instance.
- In binary relationship there are four possible constraints;
 - total-total,
 - total-partial,
 - partial-total, and
 - partial-partial.



- For a total constraint the edge leading from a relationship to an entity is drawn as a **double line**.
 - E.g., Every employee instance is part of a team instance.


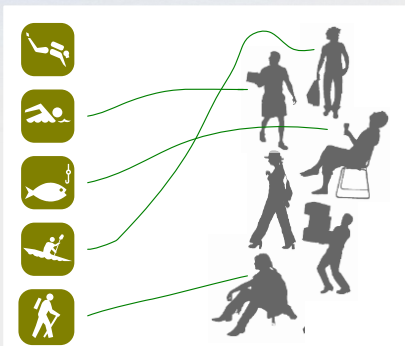
Joseph Vella - Introduction to Databases 20 Logical Database Design


20



Relationships – 1-1 Participation

- Instructor relationship between outdoor sports and employees entities.



 If a) a sport instance has to have one and only one instructor instance and b) each instructor instance can coach one and only one sport instance, then the line relationship participation constraints between **Sport** and **Instructor** each need to be total (double line).

21


Logical Database Design

Joseph Vella - Introduction to Databases

21

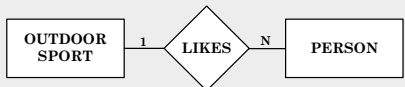
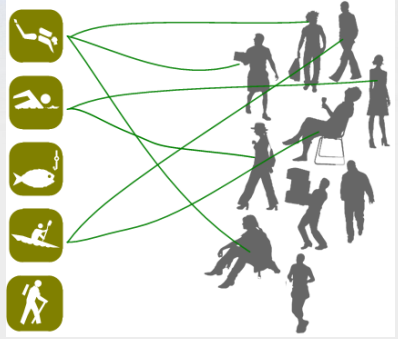
Logical Database Design

21



Relationships – 1-N Participation

- Likes relationship between outdoor sports and persons entities.



22


Logical Database Design

Joseph Vella - Introduction to Databases

22

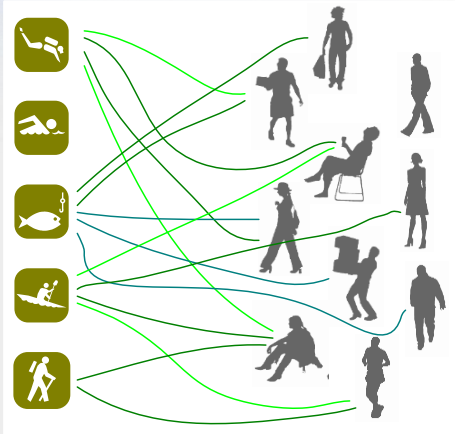

Logical Database Design

22



Relationships – N-M Participation


- Team work building relationship between outdoor sports and employees entities.

OUTDOOR SPORT — N — TEAM WORK — M — EMPLOYEE

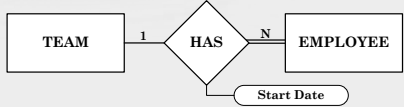
Joseph Vella - Introduction to Databases 23 Logical Database Design

23



Relationship Attributes


- Relationships might have attributes that describe them.
 - This is denoted in an ERM diagram by connecting these attribute to the respective diamond.
 - A typical example of a relationship attribute is the case of the *Starting_date* for the *Employs* relationship between the entities *TEAM* and *EMPLOYEE*.



- It is not uncommon, in a binary 1-N, that a relationship attribute is put with the other attributes of the many cardinality side entity (e.g. EMPLOYEE).
 - This is of course a "legal" arrangement but of questionable style even if the semantics of the requirements are not exactly compromised.

Joseph Vella - Introduction to Databases 24 Logical Database Design


24



Weak Entities


Joseph Vella - Introduction to Databases 25 Logical Database Design

25



A Weak entity and weak relationship


- It's not always the case that entity's instances have an independent existence from all other entity's instances!
 - Let assume that an outdoor sport requires its specialised equipment (i.e. not usable in other sport!), if we are no longer interested in trekking then all trekking equipment is no longer required!



- The ERM introduces another structure called a **weak entity** (depicted as a **double lined rectangle**) whose instances existence is related to a "normal" entity's instance existence.
- The relationship between a weak entity and its normal entity is also specialised – it's called a **weak relationship** - it's depicted as a **double lined diamond**.
- Notes:
 - A weak entity's relationship to its associated normal entity must be total.
 - We need to check primary key constraint of the weak entities instances.

Joseph Vella - Introduction to Databases 26 Logical Database Design


26



How to ERM!

Joseph Vella - Introduction to Databases 27 Logical Database Design

27




How to ERM?

- ✓ Read into subject (e.g. data & functional requirements)
- ✓ Identify entities (e.g. nouns)
 - ✓ Identify attributes
 - ✓ Identify type of attributes
 - ✓ value vs computed
 - ✓ simple vs composite
 - ✓ single vs multiple
- ✓ Identify relationships
 - ✓ Hint: see if you have 'like' attributes on entities that can be converted into relationships
- ✓ Identify weak entities & weak relationships
- ✓ Check against:
 - ✓ Requirements
 - ✓ Possible design problems ...

Joseph Vella - Introduction to Databases 28 Logical Database Design


28



Other stuff

Joseph Vella - Introduction to Databases 29 Logical Database Design

29




Advantage to ERM

- ERM are high-level languages that are relatively easy to learn by a wide spectrum of users.
- ERM are mostly independent from any particular data model and especially any database physical constructs (e.g. indices, or data placement policies).
- ERM describe a high proportion of structures, relationships and constraints information systems require (ie completeness).
 - These models are in fact used in describing the conceptual and external schemas of the ANSI/SPARC three level data architecture.
- ERM are used as a basis for building co-operative information systems (e.g. a loosely couple multi-database framework).
- ERM are also used as a test bed during analysis and early design phases.
 - The diagrams are useful to verify and validate an information system's high-level requirements.

Joseph Vella - Introduction to Databases 30 Logical Database Design

30




Problems with ERMs

- Perennial question for database designer:
Is this an entity or an entity attribute?
 - *Hint: look in requirements!?*
- Does not *scale-up*.
- For conceptual schema design there is nothing that addresses data redundancy ... designer experience has to take care of it.
- Design patterns that lead to *connection traps*:
 - Chasm trap;
 - Fan trap.
 - **Hint: Re-design.**

Joseph Vella - Introduction to Databases
31
Logical Database Design

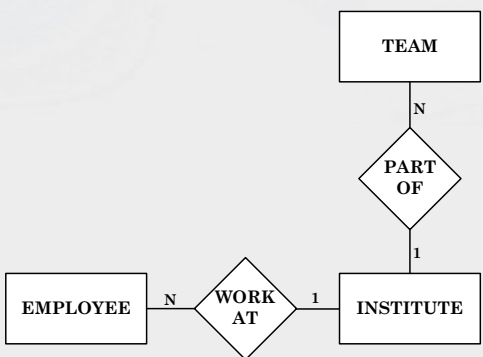
31



Connection Trap in ERM: Fan trap


Fan trap

- It is of course impossible to determine which **EMPLOYEE** instance works-on which **TEAM** instance.



Joseph Vella - Introduction to Databases
32
Logical Database Design

32




Connection Trap in ERM: Chasm trap

- Some instances of PROJECT may not connect to any TEAM instances through EMPLOYEE instances!

Joseph Vella - Introduction to Databases 33 Logical Database Design

33




A Design Pattern Issue: Are two binary relationships equal to a ternary?

- No!
As a ternary relationship can contain more relationship instances. In the lower design we cannot determine whether if a SUPPLIER instance has actually supplied a 'specific' PART to a 'specific' PROJECT.

Joseph Vella - Introduction to Databases 34 Logical Database Design

34



Case Study: Company personnel

The following description is the basis for our design:

1. The company is organised into departments. Each department has a unique name, a unique number, and a particular employee who manages the department. We keep track of the start date when that employee began managing the department. A department may have several locations.
2. A department controls a number of projects, each of which has a unique name, a unique number, and a single location.
3. We store each employee's name, national insurance number, address, salary, sex and birth date. An employee is assigned to one department but may work on several projects, which are not necessarily controlled by the same department. We keep track of the number of hours per week that an employee works on each project. We also keep track of the direct supervisor of each employee.
4. We want to keep track of the dependents of each employee for insurance purposes. We keep each dependent's name, sex, birth date, and relationship to the employee.

Joseph Vella - Introduction to Databases 35 Logical Database Design

35

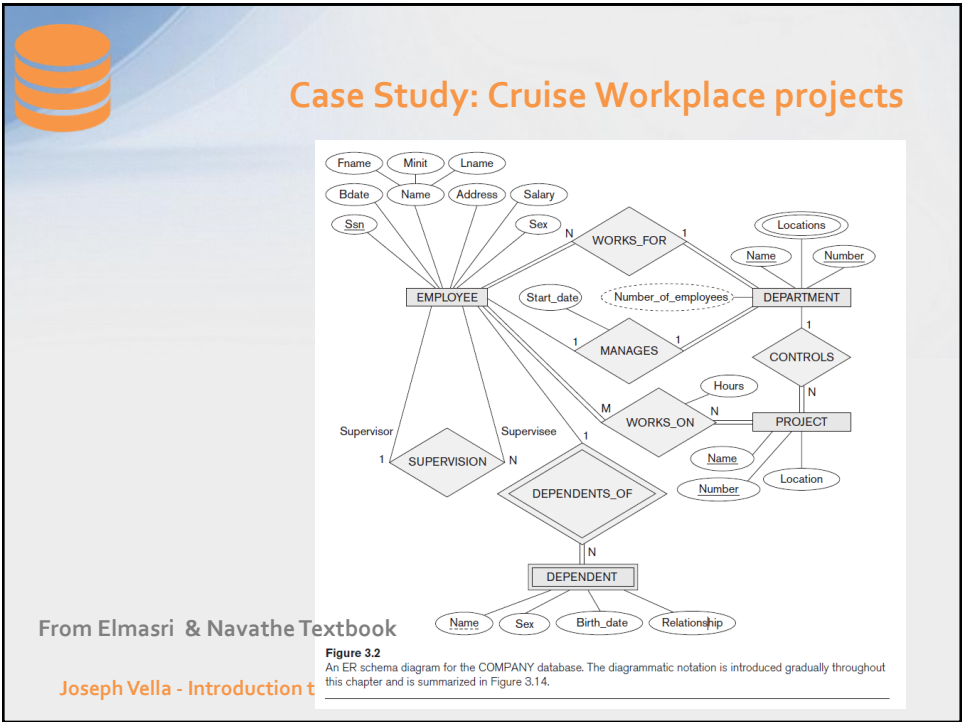


Case Study: Cruise booking

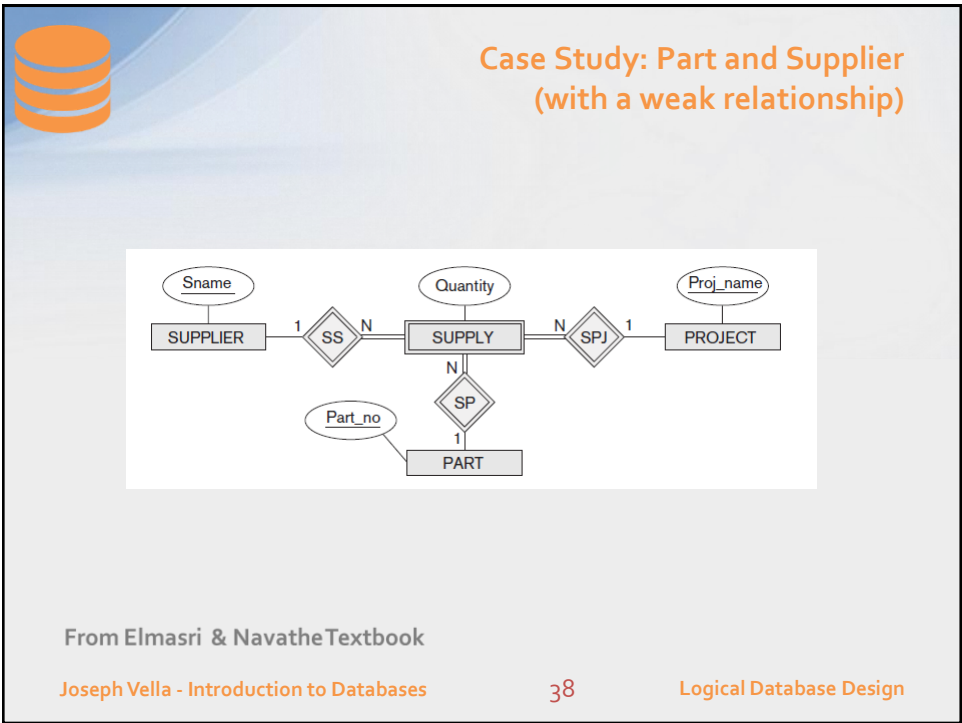
- *The following description is the basis for our design:*
- You have been entrusted to design the reservation system for Relax Cruises Co Ltd. Relax Cruises owns a number of ships each having a number of accommodation rooms in different classes. Each cruise is made by a ship leaving on a certain date between two ports in different countries. A reservation has to be made for each individual passenger and a reservation is considered to be cancelled if not paid in full two weeks before voyage commences. The Relax Cruises company has also stipulated that if the cruise a passenger wants is fully booked, the system must be capable of suggesting alternative solution either on departure days within the same week or from the closest departure or arrival ports in the same countries as originally requested.

Joseph Vella - Introduction to Databases 36 Logical Database Design


36



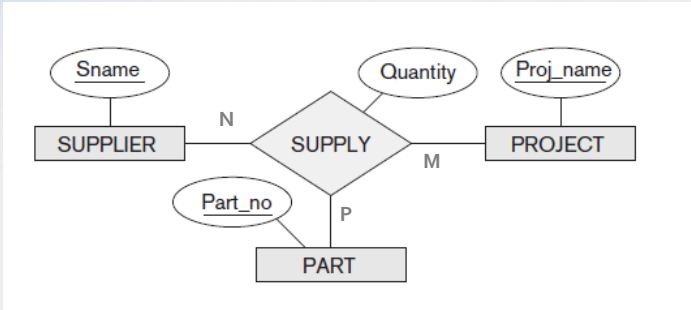
37



38




Case Study: Part and Supplier (with a ternary relationship)



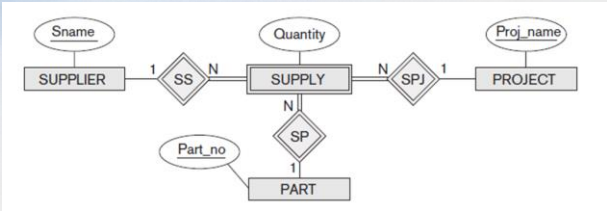
From Elmasri & Navathe Textbook

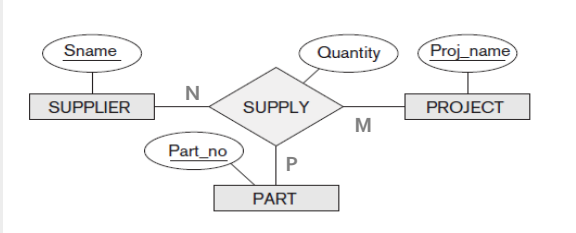
Joseph Vella - Introduction to Databases 39 Logical Database Design

39



Are these design the same?





Joseph Vella - Introduction to Databases 40 Logical Database Design

40