

Prof. Joseph G. Vella ©
joseph.g.vella@um.edu.mt

1

Quick overview of SQL

Joseph Vella - (c) - SQL - Introduction

2

Features

- SQL (Structured Query Language) is a:
 - High level;
 - Dedicated to data handling requirements;
 - Widely implemented (by different suppliers) and available over many platforms;
 - Allows connectivity to many programming languages (and other databases – SQL based and not);
 - Many SQL implementers provide procedural languages extensions that are executed by the data server process.
 - Some other actually allow coding with well known programming languages, e.g. Java and C (C++).
- SQL:
 - Some parts are declarative (others are procedural);
 - The language constructs operate over sets (actually bags) of tuples.
 - It's origin is relational theory (e.g. tables) but has long supported other structural extension to flat tuples:
 - Nested relations, XML, JSON, key-values lists etc.

Joseph Vella - (c) - SQL - Introduction

3


History

- The first SQL system is attributed to IBM (Chamberlin & Boyce):
 - It was based on CJ Codd work on relational database theory.
- Oracle, in 1979, was the first company to market & sell an SQL based system.
- The first standard came in mid eighties (SQL87) and some minor changes in 1989 (SQL89).
- A good effort was done on the second version of the standard – called SQL2. The year was 1992.
 - Still relevant with as many products still adhere to its specs.
- The next standard, called SQL3, took longer to develop, 1999, and one can say that SQL is technically a complete programming language.
 - Then followed with extensions to SQL3 – first XML & windows functions, second aligning XML with W3C, third revamp of triggers (attach to views), forth temporal extensions, fifth pattern matching and JSON.

Joseph Vella - (c) - SQL - Introduction

4

Option taken ... enough said!



Joseph Vella - (c) - SQL - Introduction

5

General Structure of simple SELECT statement is:

SELECT [DISTINCT] list of attributes
FROM list of tables
WHERE row level selection condition expression
ORDER BY list of attributes;

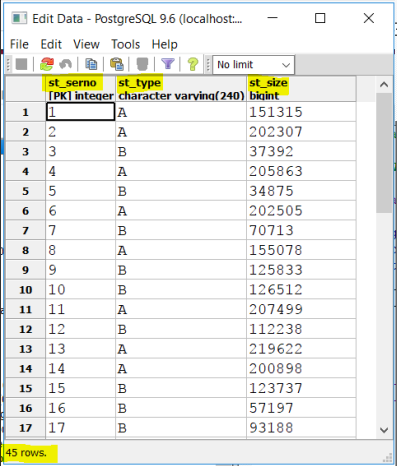
Rows Selection

Row Restrictions (aka Selection predicates)

Joseph Vella - (c) - SQL - Introduction

6

Overview of table
Walmart.stores



```
CREATE TABLE "Walmart".store
(
  st_serno integer NOT NULL,
  st_type character varying(240) NOT NULL,
  st_size bigint NOT NULL,
  CONSTRAINT st_pk PRIMARY KEY (st_serno),
  CONSTRAINT st_ck_type CHECK (st_type::text = ANY (ARRAY['A
  WITH (
    OIDS=FALSE
  );
ALTER TABLE "Walmart".store
  OWNER TO "WillStudent";
```

- How to retrieve all data in a table? (no selection!?).

```
SELECT st_serno, st_type, st_size
FROM "Walmart".store;

or

SELECT *
FROM "Walmart".store;

or

SELECT s.*
FROM "Walmart".store s;
```

Joseph Vella - (c) - SQL - Introduction

7

Simple Selection
Conjunction

- Simple selection condition based on conjunction;
 - all conjuncts must be true.

```
SELECT s.st_serno, s.st_size
FROM "Walmart".store s
WHERE st_type = 'A';
```

- Simple selection conjunction (AND operator):

```
SELECT st_serno, st_size
FROM "Walmart".store
WHERE st_type = 'A'
AND st_size < 100000;
```

- Again a simple selection with three conjuncts:

```
SELECT st_serno, st_type, st_size/100000 as megasize
FROM "Walmart".store
WHERE st_type <> 'A'
AND st_size >= 100000
AND st_serno < 25;
```

Expression with
an alias (for
column).

st_serno	st_type	megasize
integer	character varying(240)	bigint
9	B	1
10	B	1
12	B	1
15	B	1
18	B	1
21	B	1
22	B	1
23	B	1

WHERE Operator	Meaning
=	equals
<>	does not equal
>	is greater than
<	is less than
>=	is greater than or equal to
<=	is less than or equal to

Joseph Vella - (c) - SQL - Introduction

8

Simple Selection
Disjunction

- Simple selection condition based on disjunction;
 - At least one conjunct must be true.

```
SELECT st_serno, st_size
FROM "Walmart".store
WHERE st_type = 'A'
      OR st_type = 'B';
```

```
SELECT st_serno, st_size
FROM "Walmart".store
WHERE st_type = 'A'
      OR st_type = 'B'
      OR st_size > 750000;
```

Joseph Vella - (c) - SQL - Introduction

9

Conjunction & Disjunction

- The use of brackets gets useful as in the following combination of conjunction and disjunction show:

```
SELECT st_serno, st_type, st_size
FROM "Walmart".store
WHERE ( st_type = 'B'
      AND st_size < 50000 )
      OR st_size > 150000;
```

Joseph Vella - (c) - SQL - Introduction

10

Negation (as in NOT)

- Simple example of negation as in the “pattern is not found” in over table data or “pattern is not derivable” by an expression over table data.

```
SELECT st_serno, st_type, st_size
FROM "Walmart".store
WHERE NOT st_type = 'A';
```

-- one could use st_type <> 'A' too

```
SELECT st_serno, st_type, st_size
FROM "Walmart".store
WHERE NOT ( st_type = 'A'
           AND st_type = 'B' );
```

```
SELECT st_serno, st_type, st_size
FROM "Walmart".store
WHERE NOT st_type = 'A'
      AND NOT st_type = 'B';
```

De Morgan’s Laws

p	q	p and q	p or q	not (p and q)	not (p or q)
TRUE	TRUE	TRUE	TRUE	FALSE	FALSE
TRUE	FALSE	FALSE	TRUE	TRUE	FALSE
FALSE	TRUE	FALSE	TRUE	TRUE	FALSE
FALSE	FALSE	FALSE	FALSE	TRUE	TRUE

p	q	not p	not q	not p and not q	not p or not q
TRUE	TRUE	FALSE	FALSE	FALSE	FALSE
TRUE	FALSE	FALSE	TRUE	FALSE	TRUE
FALSE	TRUE	TRUE	FALSE	FALSE	TRUE
FALSE	FALSE	TRUE	TRUE	TRUE	TRUE

Joseph Vella - (c) - SQL - Introduction

11

The BETWEEN shortcut

- The BETWEEN operator abbreviates an AND expression with greater than or equal to (>=) and less than or equal to (<=) operators:

```
SELECT st_serno, st_type, st_size
FROM "Walmart".store
WHERE st_size BETWEEN 75000 AND 100000 ;
```

Joseph Vella - (c) - SQL - Introduction

12

The IN operator

- This is a very useful operator. Let's start with the syntax (and a simple example):

```
SELECT st_serno, st_type, st_size
FROM "Walmart".store
WHERE st_type IN ( 'A', 'B');
```

-- think of ('A', 'B') as a set of tuples {[attr:'A'],[attr:'B']}
- But look at this data driven query:

```
SELECT st_out.st_serno, st_out.st_type, st_out.st_size
FROM "Walmart".store AS st_out
WHERE st_out.st_type
IN (SELECT st_in.st_type
    FROM "Walmart".store AS st_in
    WHERE st_in.st_size BETWEEN 25000 AND 35000
);
```

The inner query computes "B".

	st_type character varying(240)
1	B

Joseph Vella - (c) - SQL - Introduction

13

The Nulls

- Nulls reduce our logic to three valued!?
 - E.g.
 - SELECT 1+null "result"; -> null
 - SELECT 1+coalesce(null,0) "result"; -> 1
 - COALESCE() function
 - if first argument's value is null then replace it with second argument's value

```
SELECT fe_store, fe_date, fe_temp_in_f, fe_fuel, fe_markdown1
FROM "Walmart".feature
WHERE fe_store = 1
    AND fe_markdown1 IS NULL;
```

- And how to handle it on the output:

```
SELECT fe_store, fe_date, fe_temp_in_f, fe_fuel, COALESCE (fe_markdown1,0)
FROM "Walmart".feature
WHERE fe_store = 1
    AND fe_markdown1 IS NULL; -- negation of IS NULL ->>> IS NOT NULL
```

Joseph Vella - (c) - SQL - Introduction

14

Relational Joins

Joseph Vella - (c) - SQL - Introduction

15

General Structure of simple SELECT statement is:

SELECT [DISTINCT] list of attributes

FROM list of tables

WHERE row level selection condition expression

ORDER BY list of attributes;

Joseph Vella - (c) - SQL - Introduction

16

CJ Date example schema

- CJ Date database is well known with simple structure and easy to recall it's data content (i.e. easy to verify yourself that all is well with your query attempt).
 - PLEASE note we add some tuples to the database to make some queries and their responses easier to follow.
- The database is about works scheduling and has four tables called:
 - Product (p), Supplier (s), Job (j) and Works Schedule (spj).
 - Each table has a PK and FKs are found in Works table that relates (in many to one relationship) with each of the other tables.

NOTE:

- Please run the following amendments:

```
ALTER TABLE date.p ALTER COLUMN weight
TYPE integer USING (weight::integer);
INSERT INTO date.p (p, pname, colour, weight, city)
VALUES ('P7','LOCK NUT','GREY',13,'PALO ALTO');
```

Joseph Vella - (c) - SQL - Introduction

17

Date's SPJ schema

S aka Supplier	P aka Product	J aka Job	SPJ aka Works
<u>S</u> SNAME STATUS CITY	<u>P</u> PNAME COLOUR WEIGHT CITY	<u>J</u> JNAME CITY	<u>S</u> <u>P</u> <u>J</u> QTY

S

Table

P

Primary Key Attribute

QTY

Attribute

→

Foreign Key Relationship

Joseph Vella - (c) - SQL - Introduction

18

Date's SPJ Data
(with some additions!?)

P	p	pname	colour	weight	city
	[PK]	character(20)	character(20)	integer	character(20)
1	P1	NUT	RED	12	LONDON
2	P2	BOLT	GREEN	17	PARIS
3	P3	SCREW	BLUE	17	ROME
4	P4	SCREW	RED	14	LONDON
5	P5	CAM	BLUE	12	PARIS
6	P6	COG	RED	19	LONDON
7	P7	LOCK NUT	GREY	13	PALO ALTO

S	s	sname	status	city
	[PK]	character(20)	integer	character(20)
1	S1	SMITH	20	LONDON
2	S2	JONES	10	PARIS
3	S3	BLAKE	30	PARIS
4	S4	CLARK	20	LONDON
5	S5	ADAMS	30	ATHENS

J	j	jname	city
	[PK]	character(20)	character(20)
1	J1	SORTER	PARIS
2	J2	DISPLAY	ROME
3	J3	OCR	ATHENS
4	J4	CONSOLE	ATHENS
5	J5	RAID	LONDON
6	J6	EDS	OSLO
7	J7	TAPE	LONDON

SPJ	s	p	j	qty
	[PK] ch	[PK] c	[PK] c	integer
1	S1	P1	J1	200
2	S1	P1	J4	700
3	S2	P3	J1	400
4	S2	P3	J2	200
5	S2	P3	J3	200
6	S2	P3	J4	500
7	S2	P3	J5	600
8	S2	P3	J6	400
9	S2	P3	J7	800
10	S2	P5	J2	100
11	S3	P3	J1	200
12	S3	P4	J2	500
13	S4	P6	J3	300
14	S4	P6	J7	300
15	S5	P1	J4	100
16	S5	P2	J2	200
17	S5	P2	J4	100
18	S5	P3	J4	200
19	S5	P4	J4	800
20	S5	P5	J4	400
21	S5	P5	J5	500
22	S5	P5	J7	100
23	S5	P6	J2	200
24	S5	P6	J4	500

Joseph Vella - (c) - SQL - Introduction

19

Inner Joins
based on the Equality operator

Relational Joins

Joseph Vella - (c) - SQL - Introduction

20

Equi-Join between 2 tables – using FK-PK links

Give **work schedule** details with **product's** colour and weight for RED coloured parts.

Note **wrk.p** is a FK in table **spj** and **prd.p** is a PK in table **p**.

```
SELECT wrk.s, wrk.j, prd.p, prd.colour, prd.weight
FROM date.spj wrk, date.p prd
WHERE wrk.p = prd.p
      AND prd.colour = 'RED';
```

```
SELECT wrk.s, wrk.j, prd.p, prd.colour, prd.weight
FROM date.spj wrk INNER JOIN
      date.p prd ON (wrk.p = prd.p)
WHERE prd.colour = 'RED';
```

	s	j	p	colour	weight
	charac	charac	charac	charac	integer
1	S1	J1	P1	RED	12
2	S1	J4	P1	RED	12
3	S3	J2	P4	RED	14
4	S4	J3	P6	RED	19
5	S4	J7	P6	RED	19
6	S5	J2	P6	RED	19
7	S5	J4	P1	RED	12
8	S5	J4	P4	RED	14
9	S5	J4	P6	RED	19

Joseph Vella - (c) - SQL - Introduction

21

Equi-Join between 3 tables – using FK-PK links

Give **work schedule** details with **product's** colour and weight for non RED coloured parts, and **supplier's** city.

In the first join we have **wrk.p** is a FK in table **spj** and **prd.p** is a PK in table **p**, and in the second join **spl.s** is a PK in table **s** and **wrk.s** is a FK in table **spj**

```
SELECT wrk.s, wrk.j, prd.p,
      prd.colour, prd.weight, spl.city
FROM date.spj wrk, date.p prd, date.s spl
WHERE wrk.p = prd.p
      AND wrk.s = spl.s
      AND prd.colour <> 'RED';
```


```
SELECT wrk.s, wrk.j, prd.p, prd.colour,
      prd.weight, spl.city
FROM date.spj wrk INNER JOIN
      date.p prd ON (wrk.p = prd.p)
      INNER JOIN date.s spl ON (wrk.s = spl.s)
WHERE prd.colour <> 'RED';
```

Joseph Vella - (c) - SQL - Introduction

22

Equi-Join between 3 tables
– using FK-PK links

```
SELECT wrk.s, wrk.j, prd.p, prd.colour,
      prd.weight, spl.city
FROM date.spj wrk INNER JOIN
      date.p prd ON (wrk.p = prd.p)
      INNER JOIN date.s spl ON (wrk.s = spl.s)
WHERE prd.colour <> 'RED';
```



	s	j	p	colour	weight	city
	charac	charac	charac	charac	integer	character(20)
1	S2	J1	P3	BLUE	17	PARIS
2	S2	J2	P3	BLUE	17	PARIS
3	S2	J3	P3	BLUE	17	PARIS
4	S2	J4	P3	BLUE	17	PARIS
5	S2	J5	P3	BLUE	17	PARIS
6	S2	J6	P3	BLUE	17	PARIS
7	S2	J7	P3	BLUE	17	PARIS
8	S2	J2	P5	BLUE	12	PARIS
9	S3	J1	P3	BLUE	17	PARIS
10	S5	J2	P2	GREE	17	ATHENS
11	S5	J4	P2	GREE	17	ATHENS
12	S5	J5	P5	BLUE	12	ATHENS
13	S5	J7	P5	BLUE	12	ATHENS
14	S5	J4	P3	BLUE	17	ATHENS
15	S5	J4	P5	BLUE	12	ATHENS

Joseph Vella - (c) - SQL - Introduction

23

Equi-Join between 4 tables
– using FK-PK links

```
SELECT wrk.s, wrk.j, prd.p, prd.colour, prd.weight, spl.city, job.city
FROM date.spj wrk, date.p prd, date.s spl, date.j job
WHERE wrk.p = prd.p
      AND wrk.s = spl.s
      AND wrk.j = job.j
      AND prd.colour <> 'RED';
```

```
SELECT wrk.s, wrk.j, prd.p, prd.colour, prd.weight, spl.city, job.city
FROM date.spj wrk INNER JOIN
      date.p prd ON (wrk.p = prd.p) INNER JOIN
      date.s spl ON (wrk.s = spl.s) INNER JOIN
      date.j job ON (wrk.j = job.j)
WHERE prd.colour <> 'RED';
```

Joseph Vella - (c) - SQL - Introduction


24

Equi-Join between 4 tables
– using FK-PK links

```
SELECT wrk.s, wrk.j, prd.p,
       prd.colour, prd.weight,
       spl.city, job.city

FROM date.spj wrk,
     date.p prd,
     date.s spl,
     date.j job

WHERE wrk.p = prd.p
      AND wrk.s = spl.s
      AND wrk.j = job.j
      AND prd.colour <> 'RED';
```



	s	j	p	colour	weight	city
	charac	charac	charac	charac	integer	character(20)
1	S2	J1	P3	BLUE	17	PARIS
2	S2	J2	P3	BLUE	17	PARIS
3	S2	J3	P3	BLUE	17	PARIS
4	S2	J4	P3	BLUE	17	PARIS
5	S2	J5	P3	BLUE	17	PARIS
6	S2	J6	P3	BLUE	17	PARIS
7	S2	J7	P3	BLUE	17	PARIS
8	S2	J2	P5	BLUE	12	PARIS
9	S3	J1	P3	BLUE	17	PARIS
10	S5	J2	P2	GREE	17	ATHENS
11	S5	J4	P2	GREE	17	ATHENS
12	S5	J5	P5	BLUE	12	ATHENS
13	S5	J7	P5	BLUE	12	ATHENS
14	S5	J4	P3	BLUE	17	ATHENS
15	S5	J4	P5	BLUE	12	ATHENS

Joseph Vella - (c) - SQL - Introduction

25

Equi-Join between tables
using FK-PK links & others

Give details of work schedule for RED coloured parts and where jobs and parts are collocated. In the first join we have wrk.p is a FK in table spj and prd.p is a PK in table p, in the second join spl.s is a PK in table s and wrk.s is a FK in table spj and in the third join condition we require colocation (have the same city value) between project and supplier.

```
SELECT wrk.s, wrk.j, prd.p,
       prd.colour, prd.weight, spl.city

FROM date.spj wrk,
     date.p prd,
     date.s spl

WHERE wrk.p = prd.p
      AND wrk.s = spl.s
      AND prd.city = spl.city
      AND prd.colour <> 'RED';
```

```
SELECT wrk.s, wrk.j, prd.p,
       prd.colour, prd.weight, spl.city

FROM date.spj wrk INNER JOIN
     date.p prd ON (wrk.p = prd.p) INNER JOIN
     date.s spl ON (wrk.s = spl.s AND prd.city = spl.city)

WHERE prd.colour <> 'RED';
```

	s	j	p	colour	weight	city
	charac	charac	charac	charac	integer	character(20)
1	S2	J2	P5	BLUE	12	PARIS

Joseph Vella - (c) - SQL - Introduction

26

Outer Joins

based on the Equality operator

Relational Joins

Joseph Vella - (c) - SQL - Introduction

27

Let's look at Cities values
in our tables!

Print cities values in the products and jobs tables; eliminate duplicate values and sort output by city names.

```
SELECT DISTINCT city
FROM date.p prd
ORDER BY city;
```

LONDON
PALO ALTO
PARIS
ROME

```
SELECT DISTINCT city
FROM date.j job
ORDER BY city;
```

ATHENS
LONDON
OSLO
PARIS
ROME

Print all cities values in the products and jobs tables in **one** list; eliminate duplicate values and sort output by city names.

```
SELECT DISTINCT city
FROM date.p prd
UNION
SELECT DISTINCT city
FROM date.j job
ORDER BY city;
```

ATHENS
LONDON
OSLO
PALO ALTO
PARIS
ROME

Joseph Vella - (c) - SQL - Introduction

28

Lossy Joins

Find all collocated **parts (p)** and **jobs (j)**.

```
SELECT prd.p, job.j, prd.city
FROM date.p prd INNER JOIN
      date.j job ON prd.city = job.city
ORDER BY prd.city;
```

- Cities on Output are:
LONDON
PARIS
ROME
- But we **lost!**
ATHENS
OSLO
PALO ALTO

We need
Outer Joins!

Joseph Vella - (c) - SQL - Introduction

29

Basic Right Outer Join Example

- What if we want to see all co-located **parts** and **jobs** and ensure all job cities are mentioned; even if not co-located with any part.

```
SELECT prd.p, job.j, job.city
FROM date.p prd RIGHT OUTER JOIN
      date.j job ON prd.city=job.city
ORDER BY job.city;
```

- RIGHT OUTER JOIN means keep all of the RIGHT table's (i.e. job tuples) attribute values.
- Note: the "null" values for columns that come from the non right table(s) where no match is satisfied.

	p character(20)	j character(20)	city character(20)
1		J3	ATHENS
2		J4	ATHENS
3	P1	J7	LONDON
4	P4	J5	LONDON
5	P1	J5	LONDON
6	P6	J7	LONDON
7	P4	J7	LONDON
8	P6	J5	LONDON
9		J6	OSLO
10	P2	J1	PARIS
11	P5	J1	PARIS
12	P3	J2	ROME

Joseph Vella - (c) - SQL - Introduction

30

Handling Nulls from Outer Joins

- Note: the "null" values for columns that come from the non right table(s) where no match is satisfied.
- IF you do not want to generate any nulls in our output THEN use the COALESCE() function.

```
SELECT COALESCE(prd.p, 'NO PART!') AS "prd.d",
       job.j, job.city
FROM date.p prd RIGHT OUTER JOIN
      date.j job ON prd.city=job.city
ORDER BY job.city;
```

	prd.d bpchar	j chara	city character(20)
1	NO PART!	J3	ATHENS
2	NO PART!	J4	ATHENS
3	P1	J7	LONDON
4	P4	J5	LONDON
5	P1	J5	LONDON
6	P6	J7	LONDON
7	P4	J7	LONDON
8	P6	J5	LONDON
9	NO PART!	J6	OSLO
10	P2	J1	PARIS
11	P5	J1	PARIS
12	P3	J2	ROME

Joseph Vella - (c) - SQL - Introduction

31

Basic Left Outer Join Example

- What if we want to see all co-located parts and jobs but ensure all part's cities are mentioned (even if not co-located with any job)

```
SELECT prd.p, job.j, prd.city
FROM date.p prd LEFT OUTER JOIN
      date.j job ON prd.city=job.city
ORDER BY prd.city;
```

	p chara	j chara	city character(20)
1	P1	J7	LONDON
2	P1	J5	LONDON
3	P4	J7	LONDON
4	P4	J5	LONDON
5	P6	J7	LONDON
6	P6	J5	LONDON
7	P7		PALO ALTO
8	P2	J1	PARIS
9	P5	J1	PARIS
10	P3	J2	ROME

Joseph Vella - (c) - SQL - Introduction

32

Basic Full Outer Join Example

- What if we want to see all co-located parts and jobs but ensure all parts and job cities are mentioned (even if not co-located).

```
SELECT prd.p, job.j,  
       prd.city AS "PartCity",  
       job.city AS "JobCity"  
FROM   date.p prd FULL OUTER JOIN  
       date.j job ON prd.city = job.city  
ORDER BY prd.city;
```

- FULL OUTER JOIN means keep all of the LEFT & RIGHT tables' (i.e. product and job) attribute values.

	p charac	j charac	PartCity character(20)	JobCity character(20)
1	P1	J7	LONDON	LONDON
2	P6	J5	LONDON	LONDON
3	P1	J5	LONDON	LONDON
4	P6	J7	LONDON	LONDON
5	P4	J7	LONDON	LONDON
6	P4	J5	LONDON	LONDON
7	P7	~~~~~	PALO ALTO	~~~~~
8	P5	J1	PARIS	PARIS
9	P2	J1	PARIS	PARIS
10	P3	J2	ROME	ROME
11	~~~~~	J4	~~~~~	ATHENS
12	~~~~~	J3	~~~~~	ATHENS
13	~~~~~	J6	~~~~~	OSLO

Joseph Vella - (c) - SQL - Introduction

33

Class of 2018!

Hand-drawn diagram illustrating the logic of a Full Outer Join. It shows two input tables, 'LEFT' (P) and 'RIGHT' (J), being joined. The result is a table with three rows: one for the intersection (both cities), one for the left table only (job city is null), and one for the right table only (part city is null). The diagram uses boxes, arrows, and handwritten notes like 'nulls' and '!!nulls!!' to show how data from both tables is preserved in the final output.

Joseph Vella - (c) - SQL - Introduction

34

Other Joins

(e.g. Self Joins, Non Equi-Joins)

Relational Joins

Joseph Vella - (c) - SQL - Introduction

35

Self joins are self joins characterised by using a table more than once in the from list of tables list.

An example query will be pairs of projects that are co-located.

- First attempt:

```
SELECT j1.j, j2.j, j1.city
FROM date.j j1 INNER JOIN
      date.j j2 ON( j1.city = j2.city );
```

- Second attempt i.e. solve spurious data output; we only want the two highlighted tuples as output!

```
SELECT j1.j, j2.j, j1.city
FROM date.j j1 INNER JOIN
      date.j j2 ON ( j1.city = j2.city AND j1.j > j2.j );
```

- We use these type of joins to reproduce paths and sub-structures from hierarchic structures (e.g. trees and graphs).
 - More later!?

	j charac	j charac	city character(2)
1	J1	J1	PARIS
2	J2	J2	ROME
3	J3	J4	ATHENS
4	J3	J3	ATHENS
5	J4	J4	ATHENS
6	J4	J3	ATHENS
7	J5	J7	LONDON
8	J5	J5	LONDON
9	J6	J6	OSLO
10	J7	J7	LONDON
11	J7	J5	LONDON

Joseph Vella - (c) - SQL - Introduction

36

Non-equi Joins

The join comparison expression between attributes uses a non equality operator; for example>.

Consider the query to name pairs of parts with one part "having more weight" than the second.

```
SELECT 'Part',
       p1.p,
       'has more weight than',
       p2.p,
       'by',
       p1.weight - p2.weight
FROM   date.p p1 INNER JOIN
       date.p p2 ON (p1.weight > p2.weight)
ORDER BY p1.p, p2.p;
```

	?column? unknown	p chara	?column? unknown	p char	?coli unkn	?cc int
1	Part	P2	has more weight than	P1	by	5
2	Part	P2	has more weight than	P4	by	3
3	Part	P2	has more weight than	P5	by	5
4	Part	P2	has more weight than	P7	by	4
5	Part	P3	has more weight than	P1	by	5
6	Part	P3	has more weight than	P4	by	3
7	Part	P3	has more weight than	P5	by	5
8	Part	P3	has more weight than	P7	by	4
9	Part	P4	has more weight than	P1	by	2
10	Part	P4	has more weight than	P5	by	2
11	Part	P4	has more weight than	P7	by	1
12	Part	P6	has more weight than	P1	by	7
13	Part	P6	has more weight than	P2	by	2
14	Part	P6	has more weight than	P3	by	2
15	Part	P6	has more weight than	P4	by	5
16	Part	P6	has more weight than	P5	by	7
17	Part	P6	has more weight than	P7	by	6
18	Part	P7	has more weight than	P1	by	1
19	Part	P7	has more weight than	P5	by	1

Joseph Vella - (c) - SQL - Introduction

37

Cross Product

- How can one build all possible pairs?
 - This is an example of Cross Product (or Cartesian Product).
- Output a list of all job's cities and product's colours pairs.

```
SELECT DISTINCT job.city , prd.colour
FROM   date.j job, date.p prd
ORDER BY job.city, prd.colour;
```

```
SELECT DISTINCT job.city , prd.colour
FROM   date.j job CROSS JOIN
       date.p prd
ORDER BY job.city, prd.colour;
```

- Note: Many times we do not really mean to run a product!?

	city character(2)	colour character(1)
1	ATHENS	BLUE
2	ATHENS	GREEN
3	ATHENS	GREY
4	ATHENS	RED
5	LONDON	BLUE
6	LONDON	GREEN
7	LONDON	GREY
8	LONDON	RED
9	OSLO	BLUE
10	OSLO	GREEN
11	OSLO	GREY
12	OSLO	RED
13	PARIS	BLUE
14	PARIS	GREEN
15	PARIS	GREY
16	PARIS	RED
17	ROME	BLUE
18	ROME	GREEN
19	ROME	GREY
20	ROME	RED

Joseph Vella - (c) - SQL - Introduction

38

Output Expressions

Specify out-put of data and expressions

Joseph Vella - (c) - SQL - Introduction

39

General Structure of simple SELECT statement is:

```
SELECT [DISTINCT] list of attributes  
FROM list of tables  
WHERE row level selection condition expression  
ORDER BY list of attributes;
```

Joseph Vella - (c) - SQL - Introduction

40

Simple data expressions

- One can easily choose the attributes required in the query response by listing them in **expressions** after the **SELECT** keyword (sometimes called the **SELECT list**).
 - Each expression can be an data attribute (from the tables listed in the **FROM** clause) or an expression involving the data attributes and functions available.
 - Consequently every expression in the list has a data type.
 - Also, each expression in the **SELECT** list can have an *alias*.
- We have a simple join and outputting a number of expressions (for data attributes note their data type corresponds to the data type defined in the table).
- The last expression is actually a product of two data attributes (work out the total weight of a product used in a works schedule). Note the alias specification.

```
SELECT wrk.p, wrk.s, wrk.j,  
       prd.weight, wrk.qty,  
       prd.weight*wrk.qty "total weight"  
FROM   date.p   prd INNER JOIN  
       date.spj wrk ON ( prd.p = wrk.p )  
WHERE  prd.colour = 'RED';
```

	p character(20)	s character(20)	j character(20)	weight integer	qty integer	total weight integer
1	P1	S1	J1	12	200	2400
2	P1	S1	J4	12	700	8400
3	P4	S3	J2	14	500	7000
4	P6	S4	J3	19	300	5700
5	P6	S4	J7	19	300	5700
6	P6	S5	J2	19	200	3800
7	P1	S5	J4	12	100	1200
8	P4	S5	J4	14	800	11200
9	P6	S5	J4	19	500	9500

Joseph Vella - (c) - SQL - Introduction

41

Other Expressions

- An expression in the **SELECT** list need not refer to a data attribute of a table listed in the **FROM** clause – i.e. it could be data attribute less. In this example look at **current_date**.
- One can apply a function that takes a data attribute to compute a result per data substitution.
 - In this example **initcap()** pretty prints a text string and **||** is a concatenate operator.

```
SELECT current_date "as of",  
       p, s, j, qty + 100 "new qty"  
FROM   date.spj wrk  
WHERE  wrk.qty >= 500;
```

```
SELECT jname, initcap(jname) "pretty",  
       j || ' @ ' || city "address"  
FROM   date.j job;
```

	as of date	p character(20)	s character(20)	j character(20)	new qty integer
1	2017-11-15	P1	S1	J4	800
2	2017-11-15	P3	S2	J4	600
3	2017-11-15	P3	S2	J5	700
4	2017-11-15	P3	S2	J7	900
5	2017-11-15	P4	S3	J2	600
6	2017-11-15	P5	S5	J5	600
7	2017-11-15	P4	S5	J4	900
8	2017-11-15	P6	S5	J4	600

	jname character(20)	pretty text	address text
1	SORTER	Sorter	J1 @ PARIS
2	DISPLAY	Display	J2 @ ROME
3	OCR	Ocr	J3 @ ATHENS
4	CONSOLE	Console	J4 @ ATHENS
5	RAID	Raid	J5 @ LONDON
6	EDS	Eds	J6 @ OSLO
7	TAPE	Tape	J7 @ LONDON

Joseph Vella - (c) - SQL - Introduction

42

Output Duplicate Elimination

- Consider the query: Provide part numbers for RED parts that are on a work schedule.
 - Note the duplicates in the output!?

```
SELECT wrk.p
FROM date.p prd INNER JOIN
      date.spj wrk ON ( prd.p = wrk.p )
WHERE prd.colour = 'RED';
```

- To remove any duplicates add the DISTINCT keyword just after the SELECT keyword.

```
SELECT DISTINCT wrk.p
FROM date.p prd INNER JOIN
      date.spj wrk ON ( prd.p = wrk.p )
WHERE prd.colour = 'RED';
```

	p	character(20)
1	P1	
2	P1	—
3	P4	
4	P6	
5	P6	—
6	P6	—
7	P1	—
8	P4	—
9	P6	—

	p	character(20)
1	P1	
2	P4	
3	P6	

Joseph Vella - (c) - SQL - Introduction

43

Handling Output (as in total)

Like sorting, ranking, top most, and slice.

Joseph Vella - (c) - SQL - Introduction

44

General Structure of simple SELECT statement is:

```
SELECT [DISTINCT] list of attributes
FROM list of tables
WHERE row level selection condition expression
ORDER BY list of attributes;
```

Joseph Vella - (c) - SQL - Introduction

45

Sorting the output

- Sorting is value based (e.g. on data attributes and expressions) – use the ORDER BY clause.
 - One can sort on a sequence of attributes/expressions.
 - Each expression can be sorted in ascending or descending order.
 - The default is ascending order.

- Examples:

```
SELECT j, s, p, qty + 100 "new qty"
FROM date.spj wrk
WHERE wrk.qty >= 500
ORDER BY j ASC, s ASC, p DESC;
```

```
SELECT j, s, p, qty + 100 "new qty"
FROM date.spj wrk
WHERE wrk.qty >= 500
ORDER BY j, "new qty" DESC;
```

```
SELECT s, p
FROM date.spj wrk
WHERE wrk.qty >= 500
ORDER BY qty + 100 DESC;
```

Joseph Vella - (c) - SQL - Introduction

46

Non-value based restrictions

```
SELECT select_List
  FROM table_expression
  [ ORDER BY ... ]
  [ LIMIT { number | ALL } ] [ OFFSET number ]
```

- What if we want a sub-set of the output and the subset comprehension is not value based (e.g. cannot use the WHERE clause)? One technique with modern SQL is to use ordinal directives LIMIT and OFFSET.
- With LIMIT count one specifies the maximum number of rows to output.
- With OFFSET count one specifies the rows to ignore from the start of the "original" output.

! In most cases LIMIT and OFFSET are specified with an ORDER BY expression.

Otherwise output, on the same database state, cannot be guaranteed to be the same!?

OFFSET 3

LIMIT 4

1	2	3	4	5	6	7	...
---	---	---	---	---	---	---	-----

Joseph Vella - (c) - SQL - Introduction

47

LIMIT the number of rows

```
SELECT j, s, p, qty
  FROM date.spj wrk
 WHERE wrk.qty >= 100
 ORDER BY qty DESC;
```

	j	s	p	qty
	charac	chara	charc	integer
1	J7	S2	P3	800
2	J4	S5	P4	800
3	J4	S1	P1	700
4	J5	S2	P3	600
5	J5	S5	P5	500
6	J4	S5	P6	500
7	J4	S2	P3	500
8	J2	S3	P4	500
9	J4	S5	P5	400
10	J1	S2	P3	400
11	J6	S2	P3	400
12	J3	S4	P6	300
13	J7	S4	P6	300
14	J1	S1	P1	200
15	J2	S5	P2	200
16	J2	S5	P6	200
17	J4	S5	P3	200
18	J3	S2	P3	200
19	J2	S2	P3	200
20	J1	S3	P3	200
21	J4	S5	P2	100
22	J4	S5	P1	100
23	J7	S5	P5	100
24	J2	S2	P5	100

```
SELECT j, s, p, qty
  FROM date.spj wrk
 WHERE wrk.qty >= 100
 ORDER BY qty DESC
 LIMIT 5 ;
```

	j	s	p	qty
	charac	chara	charc	integer
1	J4	S5	P4	800
2	J7	S2	P3	800
3	J4	S1	P1	700
4	J5	S2	P3	600
5	J2	S3	P4	500

Joseph Vella - (c) - SQL - Introduction

48

OFFSET – re-start of output

```
SELECT j, s, p, qty
FROM date.spj wrk
WHERE wrk.qty >= 100
ORDER BY qty DESC;
```

	j	s	p	qty
	charac	chara	charc	integer
1	J7	S2	P3	800
2	J4	S5	P4	800
3	J4	S1	P1	700
4	J5	S2	P3	600
5	J5	S5	P5	500
6	J4	S5	P6	500
7	J4	S2	P3	500
8	J2	S3	P4	500
9	J4	S5	P5	400
10	J1	S2	P3	400
11	J6	S2	P3	400
12	J3	S4	P6	300
13	J7	S4	P6	300
14	J1	S1	P1	200
15	J2	S5	P2	200
16	J2	S5	P6	200
17	J4	S5	P3	200
18	J3	S2	P3	200
19	J2	S2	P3	200
20	J1	S3	P3	200
21	J4	S5	P2	100
22	J4	S5	P1	100
23	J7	S5	P5	100
24	J2	S2	P5	100

```
SELECT j, s, p, qty
FROM date.spj wrk
WHERE wrk.qty >= 100
ORDER BY qty DESC
OFFSET 20 ;
```

	j	s	p	qty
	charac	chara	charc	integer
1	J4	S5	P2	100
2	J4	S5	P1	100
3	J7	S5	P5	100
4	J2	S2	P5	100

Joseph Vella - (c) - SQL - Introduction

49

LIMIT and OFFSET

```
SELECT j, s, p, qty
FROM date.spj wrk
WHERE wrk.qty >= 100
ORDER BY qty DESC;
```

	j	s	p	qty
	charac	chara	charc	integer
1	J7	S2	P3	800
2	J4	S5	P4	800
3	J4	S1	P1	700
4	J5	S2	P3	600
5	J5	S5	P5	500
6	J4	S5	P6	500
7	J4	S2	P3	500
8	J2	S3	P4	500
9	J4	S5	P5	400
10	J1	S2	P3	400
11	J6	S2	P3	400
12	J3	S4	P6	300
13	J7	S4	P6	300
14	J1	S1	P1	200
15	J2	S5	P2	200
16	J2	S5	P6	200
17	J4	S5	P3	200
18	J3	S2	P3	200
19	J2	S2	P3	200
20	J1	S3	P3	200
21	J4	S5	P2	100
22	J4	S5	P1	100
23	J7	S5	P5	100
24	J2	S2	P5	100

```
SELECT j, s, p, qty
FROM date.spj wrk
WHERE wrk.qty >= 100
ORDER BY qty DESC
LIMIT 5
OFFSET 10 ;
```

	j	s	p	qty
	charac	chara	charc	integer
1	J6	S2	P3	400
2	J3	S4	P6	300
3	J7	S4	P6	300
4	J1	S1	P1	200
5	J2	S5	P2	200

Joseph Vella - (c) - SQL - Introduction

50

N th Row (e.g. get the 10th)

```
SELECT j, s, p, qty
FROM date.spj wrk
WHERE wrk.qty >= 100
ORDER BY qty DESC;
```

	j	s	p	qty
	charac	chara	chara	integer
1	J7	S2	P3	800
2	J4	S5	P4	800
3	J4	S1	P1	700
4	J5	S2	P3	600
5	J5	S5	P5	500
6	J4	S5	P6	500
7	J4	S2	P3	500
8	J2	S3	P4	500
9	J4	S5	P5	400
10	J1	S2	P3	400
11	J6	S2	P3	400
12	J3	S4	P6	300
13	J7	S4	P6	300
14	J1	S1	P1	200
15	J2	S5	P2	200
16	J2	S5	P6	200
17	J4	S5	P3	200
18	J3	S2	P3	200
19	J2	S2	P3	200
20	J1	S3	P3	200
21	J4	S5	P2	100
22	J4	S5	P1	100
23	J7	S5	P5	100
24	J2	S2	P5	100

```
SELECT j, s, p, qty
FROM date.spj wrk
WHERE wrk.qty >= 100
ORDER BY qty DESC
LIMIT 1
OFFSET 10 ; -- N=10
```

	j	s	p	qty
	charac	chara	chara	integer
1	J6	S2	P3	400

Joseph Vella - (c) - SQL - Introduction

51

Functions for Output Expressions

Actually these functions, with restrictions, can be used in for example WHERE and ORDER BY expressions.

Joseph Vella - (c) - SQL - Introduction

52

Row

vs

Aggregate

vs

Hybrid (Group By queries)

vs

Data frames (e.g. windows)*

And some other stuff like adding conditional expression to SELECT list and generating data on the fly.

*) topic in forthcoming units.

Row function example

```
SELECT j, s, p, qty
FROM date.spj wrk
WHERE wrk.qty >= 500;
```

	j	s	p	qty
	charac	chara	charz	integer
1	J4	S1	P1	700
2	J4	S2	P3	500
3	J5	S2	P3	600
4	J7	S2	P3	800
5	J2	S3	P4	500
6	J5	S5	P5	500
7	J4	S5	P4	800
8	J4	S5	P6	500

```
SELECT j, s, p,
       | (power(qty,1.5))/1000 "whatever"
FROM date.spj wrk
WHERE wrk.qty >= 500;
```

	j	s	p	whatever
	charac	chara	charz	numeric
1	J4	S1	P1	18.5202591774521340
2	J4	S2	P3	11.1803398874989480
3	J5	S2	P3	14.6969384566990690
4	J7	S2	P3	22.6274169979695210
5	J2	S3	P4	11.1803398874989480
6	J5	S5	P5	11.1803398874989480
7	J4	S5	P4	22.6274169979695210
8	J4	S5	P6	11.1803398874989480

- Function power() is a row function because it output's one value for each application per row.
 - If the output has 8 rows then applying the same filter and a row function on the output one remains with 8.

Aggregate function example

```
SELECT j, s, p, qty
FROM date.spj wrk
WHERE wrk.qty >= 500;
```

	j	s	p	qty
	charac	chara	chara	integer
1	J4	S1	P1	700
2	J4	S2	P3	500
3	J5	S2	P3	600
4	J7	S2	P3	800
5	J2	S3	P4	500
6	J5	S5	P5	500
7	J4	S5	P4	800
8	J4	S5	P6	500

```
SELECT sum(qty) "all together"
FROM date.spj wrk
WHERE wrk.qty >= 500;
```

	all together
	bigint
1	4900

- Function sum() is an aggregate function because it output's one value for all rows on the output.
 - If the output has 8 rows then applying the same filter and an aggregate function on the output reduces to one row.

Joseph Vella - (c) - SQL - Introduction

55

Aggregate functions example

```
SELECT j, s, p, qty
FROM date.spj wrk
WHERE wrk.qty >= 500;
```

	j	s	p	qty
	charac	chara	chara	integer
1	J4	S1	P1	700
2	J4	S2	P3	500
3	J5	S2	P3	600
4	J7	S2	P3	800
5	J2	S3	P4	500
6	J5	S5	P5	500
7	J4	S5	P4	800
8	J4	S5	P6	500

One can have **multiple aggregate functions** in one output list

```
SELECT sum(qty) "Sum",
       max(qty) "Max",
       count(qty) "N"
FROM date.spj wrk
WHERE wrk.qty >= 500;
```

	Sum	Max	N
	bigint	integer	bigint
1	4900	800	8

But, what is happening here?

```
SELECT 'sum' "agg_function", sum(qty) "value"
FROM date.spj wrk
WHERE wrk.qty >= 500
UNION ALL
SELECT 'count', count(qty)
FROM date.spj wrk
WHERE wrk.qty >= 500
UNION ALL
SELECT 'max', max(qty)
FROM date.spj wrk
WHERE wrk.qty >= 500;
```


	agg_function	value
	text	bigint
1	max	800
2	count	8
3	sum	4900

Joseph Vella - (c) - SQL - Introduction

56

Cannot mix row & aggregate functions ...

- We cannot mix row and aggregate functions in our SELECT ... FROM ... WHERE syntax.
- To do that we need additional syntax!? For example ,the GROUP BY syntax.



```
SELECT j, s, p,  
       (power(qty,1.5))/1000 "whatever",  
       sum(qty) "all together"  
FROM date.spj wrk  
WHERE wrk.qty >= 500;
```

ERROR: column "wrk.j" must appear
in the GROUP BY clause or be used in
an aggregate function
LINE 1: SELECT j, s, p,

Joseph Vella - (c) - SQL - Introduction

57

Row & Aggregate Function

```
SELECT j, s, p, qty  
FROM date.spj wrk  
WHERE wrk.qty >= 250  
ORDER BY j ASC, s ASC, p DESC;
```

	j	s	p	qty
	charac	chara	charz	integer
1	J1	S2	P3	400
2	J2	S3	P4	500
3	J3	S4	P6	300
4	J4	S1	P1	700
5	J4	S2	P3	500
6	J4	S5	P6	500
7	J4	S5	P5	400
8	J4	S5	P4	800
9	J5	S2	P3	600
10	J5	S5	P5	500
11	J6	S2	P3	400
12	J7	S2	P3	800
13	J7	S4	P6	300

```
SELECT j, sum( qty )  
FROM date.spj wrk  
WHERE wrk.qty >= 250  
GROUP BY j  
ORDER BY j;
```

	j	sum
	charac	bigint
1	J1	400
2	J2	500
3	J3	300
4	J4	2900
5	J5	1100
6	J6	400
7	J7	1100

- To mix row (as in expression *j*) and aggregate (as in expression *sum(qty)*) one needs the appropriate SQL statement structure.
- The simplest is the GROUP BY construct – more later.

Joseph Vella - (c) - SQL - Introduction

58

Partitioning & Running Totals

- In normal `SELECT ... FROM ... WHERE` statements we have made it a point to say that each row, from the data sources, is evaluated independently of any other row.
 - Also the output for a row that passed the filter can only access the data in that row!
- To circumvent this we need a new syntax! For example:

```
SELECT j, p, qty, sum(qty)
      OVER (ORDER BY j)
FROM   date.spj wrk
ORDER BY j, p;
```
- The idea here is to compute a running total. The partition, or window, in this case is each distinct Job reference. For the first case we have 'J1'. All the parts allocated to 'J1', i.e. 'P1' twice and 'P3', total 800 (200+ 400 + 200).
 - The output include The Project, parts used and their respective quantity.
- The simplest is the **Windows Functions** construct – more later.

	j	p	qty	sum
	charac	character(20)	integer	bigint
1	J1	P1	200	800
2	J1	P3	400	800
3	J1	P3	200	800
4	J2	P2	200	2000
5	J2	P3	200	2000
6	J2	P4	500	2000
7	J2	P5	100	2000
8	J2	P6	200	2000
9	J3	P3	200	2500
10	J3	P6	300	2500
11	J4	P1	100	5800
12	J4	P1	700	5800
13	J4	P2	100	5800
14	J4	P3	200	5800
15	J4	P3	500	5800
16	J4	P4	800	5800
17	J4	P5	400	5800
18	J4	P6	500	5800
19	J5	P3	600	6900
20	J5	P5	500	6900
21	J6	P3	400	7300
22	J7	P3	800	8500
23	J7	P5	100	8500
24	J7	P6	300	8500

Joseph Vella - (c) - SQL - Introduction

59

Examples: numeric functions

- Note we are using the most basic syntax for exposition.

```
-- absolute value
SELECT abs(-5);
-- 5::int

-- nearest integer greater than
-- or equal to argument
SELECT ceil(-1.713);
-- -1::numeric

-- nearest integer less than
-- or equal to argument
SELECT floor(-1.713);
-- -2::numeric

-- truncates toward zero or decimal places
SELECT trunc(1.7149), trunc(1.7149,3);
-- 1::numeric, 1.714::numeric

-- round to nearest integer
SELECT round(1.7149), round(1.7149,3);
-- 2::numeric, 1.715::numeric

-- exponential & natural log
SELECT exp(2);
-- 7.389::dp

SELECT ln(7.38905609893065);
-- 2.0::numeric

-- raise to the power
SELECT power(2,3);
-- 8::dp

-- bucketting (exp, min, max, numb buckets)
SELECT width_bucket(6.5,0,10,10);
-- 7::integer

-- random number (and re-seed if required)
SELECT random();
-- 0 <= x <= 1

-- trigonometric and reverse functions
SELECT sind(30);

SELECT asind(.5);
```

Joseph Vella - (c) - SQL - Introduction

60

Example: string functions

```
-- character length
SELECT char_length('Test');
-- 4::integer

-- overlay
SELECT overlay('ABCDEF'
  placing '12'
  from 2
  for 4);
-- A12F::text

-- substring
-- (also left() and right() exist)
SELECT substring('ABCDEF' from 2 for 2);
-- BC::text

-- trim
SELECT trim(leading '*' from '*123*'),
       trim(both '*' from '*123*'),
       trim(trailing '*' from '*123*');
-- 123*::text, 123::text, *123::text

-- position
SELECT position('E' in 'ABCDEF');
-- 5::integer

-- concatenate
SELECT concat('A', 'B', 'C'),
       concat_ws('|', 'A', 'B', 'C');
-- ABC::text, A|B|C::text

-- format
SELECT format('Dear %s %s %s %s, Wellcome to ...',
  'Ms', 'Hanna', null, 'Staples');
-- Dear Ms Hanna Staples, Wellcome to ....::text
```

Joseph Vella - (c) - SQL - Introduction

61

Example: date functions

```
SELECT to_char(current_date, 'YYYY Mon DD'),
       to_char(current_date, 'J'),
       to_char(localtime, 'HH12:MI:SS');
-- "2017 Nov 29"::text, "2458087"::text, "01:12:59"::text
-- note J is for Julian Day (days since November 24, 4714 BC at midnight)

SELECT current_date + 7 "todayweek",
       date '2017-01-01' + integer '100' "100days",
       48*3600 * interval '1 second' "who_much?";
-- "2017-12-06"::date, "2017-04-11"::date, "48:00:00"::interval

-- "age" (in years and months)
SELECT age(current_date, date '2000-01-01')
-- "17 years 10 mons 28 days"::interval

-- date_part extraction
SELECT date_part('day', current_date),
       date_part('month', current_date),
       date_part('year', current_date);
-- 29::double pre, 11::double pre, 2017::double pre

-- date overlaps
SELECT ( date '2017-06-01', date '2017-07-31')
       OVERLAPS
       ( date '2017-01-01', date '2017-12-31');
-- t::boolean
```

Joseph Vella - (c) - SQL - Introduction

62

Example Aggregate Funct.s

```
-- count(*) and count( expr )
SELECT count(*),
       count(distinct s),
       count(distinct p),
       count(distinct j)
FROM date.spj wrk;
-- 24::bigint, 5::bigint, 6::bigint, 7::bigint

-- sum
SELECT sum(qty)
FROM date.spj wrk;
-- 8500::bigint

-- min & max
SELECT sum(qty), min(qty), max(qty)
FROM date.spj wrk;
-- 8500::bigint, 100::bigint, 800::bigint

-- array_agg -> get values into an array
SELECT array_agg(trim( city ))
FROM date.j job;
-- {PARIS,ROME,ATHENS,ATHENS,LONDON,OSLO,LONDON}::text[]

SELECT array_agg(trim( t.city ))
FROM (SELECT DISTINCT city FROM date.j job) AS t;
-- {OSLO,ROME,PARIS,ATHENS,LONDON}::text[]
```

Joseph Vella - (c) - SQL - Introduction

63

Quick Aside:
A set generating function!

- The function `generate_series()` returns a number of tuples on the fly. The function takes two to three arguments: first term, last term, and an optional jump.
 - Each argument is a numeric and usually is an integer.
- The `generate_series()` can be used in SELECT list expressions or as a data source in the FROM list of expressions.

```
SELECT generate_series(1,5);
SELECT generate_series(-3,1);
SELECT generate_series(1,10,2);
SELECT current_date+i.next
FROM generate_series(1,5) as i(next);
```

generate_series integer
1
2
3
4
5

generate_series integer
-3
-2
-1
0
1

generate_series integer
1
3
5
7
9

?column? date
2017-11-29
2017-11-30
2017-12-01
2017-12-02
2017-12-03

Joseph Vella - (c) - SQL - Introduction

64

More example Aggregate Functions

- The following Boolean functions enable comparisons of a value to a set.
- ALL** – all tuples have the stated property. (I.e. tuples [1], [2], [3], [4] & [5] are all greater than constant o (zero)).
- ANY** (or **SOME**) – at least one tuple has the stated property.
- EXISTS** is a test for non-empty set.
- IN** confirms presence (already covered earlier).

```
SELECT 0 <= ALL (SELECT * FROM generate_series(1,5));
-- t:Boolean

SELECT 0 <= ALL (SELECT * FROM generate_series(-3,1));
-- f:Boolean

SELECT 0 <= ANY (SELECT * FROM generate_series(1,5));
SELECT 0 <= SOME (SELECT * FROM generate_series(1,5));

-- t:Boolean
SELECT 0 <= ANY (SELECT * FROM generate_series(-3,1));
-- t:Boolean

SELECT EXISTS (SELECT * FROM generate_series(1,5));
-- t:Boolean
SELECT EXISTS (SELECT * FROM generate_series(-3,1));
-- t:Boolean
SELECT EXISTS (SELECT * FROM generate_series(-3,1) as t(a)
WHERE a > 10);
-- f:Boolean

SELECT 3 IN (SELECT * FROM generate_series(1,5));
-- t:Boolean
```

Joseph Vella - (c) - SQL - Introduction

65

Example: output conditionals

- The **CASE** expression can alter the output of a SELECT expression.
- The condition, expressed in the WHEN must return a Boolean (i.e. True or False). If the expression returns a 'True' then the expression listed after the relative THEN keyword is executed.
- Only one (the first) WHEN condition is executed.
- If no WHEN condition is true and if an ELSE expression is present then it is executed!
- The data type of all result expression must match!

```
CASE WHEN condition THEN result
[WHEN ...]
[ELSE result]
END

CASE expression
WHEN value THEN result
[WHEN ...]
[ELSE result]
END
```

```
SELECT sname,
CASE WHEN status = 10
AND status < 10 THEN 'Green'
WHEN status = 20 THEN 'Orange'
ELSE 'Red'
END -- case
FROM date.s;
```

s	sname	status	city
[PK]	character(20)	integer	character(20)
S1	SMITH	20	LONDON
S2	JONES	10	PARIS
S3	BLAKE	30	PARIS
S4	CLARK	20	LONDON
S5	ADAMS	30	ATHENS

sname	case
character(20)	text
SMITH	Orange
JONES	Red
BLAKE	Red
CLARK	Orange
ADAMS	Red

Joseph Vella - (c) - SQL - Introduction

66

Example: output conditionals

- Another two output expressions are: **COALESCE()**, seen before, and related **NULLIF()**.

```
SELECT p,
       coalesce(substring(pname from '%#"NUT#"' for '#'), 'Not nutty')
FROM date.p prd;
```

-- searching with regular expressions:
Meta character sequence # " is the return string indicator.
Therefore the above returns 'NUT' if 'NUT' is present in the string.


p	coalesce charac text
P1	NUT
P2	Not nutty
P3	Not nutty
P4	Not nutty
P5	Not nutty
P6	Not nutty
P7	NUT

```
-- nullif -- we want to generate nuls!!! whimsical i know!?
SELECT p,
       substring(pname from '%#"NUT#"' for '#'),
       nullif(substring(pname from '%#"NUT#"' for '#'), 'NUT')
FROM date.p prd
WHERE substring(pname from '%#"NUT#"' for '#') IS NOT NULL;
```

p	substring charac text	nullif text
P1	NUT	—
P7	NUT	—

```
-- greatest and least
-- (not very portable)
SELECT greatest(1,2,3,4,5), least(1,2,3,4,5);
```

greatest integer	least integer
5	1



Group By Queries

Great for cross tabulations - basic ones for now

Motivation

p	s	j	qty
charac	charac	charac	integer
P1	S1	J1	200
P1	S1	J4	700
P1	S5	J4	100
P2	S5	J2	200
P2	S5	J4	100
P3	S2	J1	400
P3	S2	J2	200
P3	S2	J3	200
P3	S2	J4	500
P3	S2	J5	600
P3	S2	J6	400
P3	S2	J7	800
P3	S3	J1	200
P3	S5	J4	200
P4	S3	J2	500
P4	S5	J4	800
P5	S2	J2	100
P5	S5	J4	400
P5	S5	J5	500
P5	S5	J7	100
P6	S4	J3	300
P6	S4	J7	300
P6	S5	J2	200
P6	S5	J4	500

- How many parts are required in all our project's work schedules?

```
SELECT sum(qty)
FROM date.spj
WHERE p='P1';
```

1000

- What if we want total for part 'P1' and 'P2'?

```
SELECT 'P1',sum(qty)
FROM date.spj
WHERE p='P1'
UNION
SELECT 'P2',sum(qty)
FROM date.spj
WHERE p='P2';
```

P1	1000
P2	300
- This gets unyielding!?
 - No idea what parts are mentioned in a works schedule.
 - Cross tab of, say p and s, gets to be a long winded process.

Joseph Vella - (c) - SQL - Introduction

69

Motivation

```
SELECT DISTINCT p
FROM date.spj
ORDER BY p;
```

```
SELECT p, qty
FROM date.spj
ORDER BY p;
```

```
SELECT DISTINCT s
FROM date.spj
ORDER BY s;
```

```
SELECT s, qty
FROM date.spj
ORDER BY s;
```

s	qty
char.	integer
S1	200
S1	700
S2	400
S2	200
S2	200
S2	500
S2	600
S2	400
S2	800
S2	100
S3	200
S3	500
S4	300
S4	300
S5	200
S5	100
S5	500
S5	100
S5	200
S5	100
S5	200
S5	800
S5	400
S5	500

p	qty
charac	integer
P1	100
P1	200
P1	700
P2	100
P2	200
P3	800
P3	200
P3	200
P3	500
P3	600
P3	400
P3	400
P3	200
P4	500
P4	800
P5	100
P5	400
P5	500
P5	100
P6	500
P6	300
P6	300
P6	200

Joseph Vella - (c) - SQL - Introduction

70

Issue!? (already mentioned)

- Remember we cannot mix row and aggregate function in a SELECT list of expressions in our simple syntax of SELECT FROM WHERE:



```
SELECT wrk.p,  
       sum(wrk.qty) "Parts"  
FROM date.spj wrk;
```

ERROR: column "wrk.p"
must appear in the GROUP
BY clause or be used in an
aggregate function

Joseph Vella - (c) - SQL - Introduction

71

The Generic GROUP BY syntax

General Structure of simple SELECT statement is:

```
SELECT [DISTINCT] list of attributes  
FROM list of tables  
WHERE row level selection condition expression  
ORDER BY list of attributes;
```

General Structure of SELECT statement with Group By

```
SELECT [DISTINCT] list of expressions  
FROM list of tables  
WHERE row level selection condition expression  
[GROUP BY aggregate on expression list  
[HAVING aggregate condition expression ]]  
ORDER BY list of attributes;
```

Joseph Vella - (c) - SQL - Introduction

72

Motivation

s	charac
S1	
S2	
S3	
S4	
S5	

s	qty	charac	integer
S1	200		
S1	700		
S2	400		
S2	200		
S2	200		
S2	500		
S2	600		
S2	400		
S2	800		
S2	100		
S3	200		
S3	500		
S4	300		
S4	300		
S5	200		
S5	100		
S5	500		
S5	100		
S5	200		
S5	100		
S5	200		
S5	800		
S5	400		
S5	500		

s	Supplier	charac	bigint
S1	900		
S2	3200		
S3	700		
S4	600		
S5	3100		

```
SELECT wrk.s,
       sum(wrk.qty)
FROM   date.spj wrk
GROUP BY wrk.s
ORDER BY wrk.s;
```

p	charac	integer
P1	100	
P1	200	
P1	700	
P2	100	
P2	200	
P3	800	
P3	200	
P3	200	
P3	500	
P3	600	
P3	400	
P3	400	
P3	200	
P4	800	
P5	100	
P5	400	
P5	500	
P5	100	
P6	500	
P6	300	
P6	300	
P6	200	

p	charac	Parts	bigint
P1	1000		
P2	300		
P3	3500		
P4	1300		
P5	1100		
P6	1300		

Row expression

Aggregate expression

"Supplier"

"Parts"

Joseph Vella - (c) - SQL - Introduction

73

Cross tabulation

```
SELECT wrk.p,
       wrk.s,
       sum(wrk.qty) "Parts",
       count(*)
FROM   date.spj wrk
GROUP BY wrk.p, wrk.s
ORDER BY wrk.p, wrk.s;
```

p	charac	s	charac	Parts	bigint	count	bigint
P1		S1		900		2	
P1		S5		100		1	
P2		S5		300		2	
P3		S2		3100		7	
P3		S3		200		1	
P3		S5		200		1	
P4		S3		500		1	
P4		S5		800		1	
P5		S2		100		1	
P5		S5		1000		3	
P6		S4		600		2	
P6		S5		700		2	

```
SELECT wrk.p,
       wrk.s,
       wrk.j,
       sum(wrk.qty) "Parts",
       count(*)
FROM   date.spj wrk
GROUP BY wrk.p, wrk.s, wrk.j
ORDER BY wrk.p, wrk.s;
```

Since p, s and j are PK then each group is a tuple!

Good idea to copy the GROUP BY expr.s in SELECT list!

Joseph Vella - (c) - SQL - Introduction

74

Using a row filter

```
SELECT 'P1', sum(qty)
FROM date.spj
WHERE p='P1'
UNION
SELECT 'P2', sum(qty)
FROM date.spj
WHERE p='P2'
```

P1	1000
P2	300

```
SELECT wrk.p, sum(wrk.qty) "Parts"
FROM date.spj wrk
WHERE wrk.p IN ('P1','P2')
GROUP BY wrk.p
ORDER BY wrk.p;
```

Joseph Vella - (c) - SQL - Introduction

75

Using a group filter (aggregate)

```
SELECT wrk.p,
       wrk.s,
       sum(wrk.qty) "Parts",
       count(*)
FROM date.spj wrk
GROUP BY wrk.p, wrk.s
ORDER BY wrk.p, wrk.s;
```

p	s	Parts	count
charac	charac	bigint	bigint
P1	S1	900	2
P1	S5	100	1
P2	S5	300	2
P3	S2	3100	7
P3	S3	200	1
P3	S5	200	1
P4	S3	500	1
P4	S5	800	1
P5	S2	100	1
P5	S5	1000	3
P6	S4	600	2
P6	S5	700	2

Take off heavy hitters!

```
SELECT wrk.p,
       wrk.s,
       sum(wrk.qty) "Parts",
       count(*)
FROM date.spj wrk
GROUP BY wrk.p,
         wrk.s
HAVING count(*) < 3
ORDER BY wrk.p, wrk.s;
```

p	s	Parts	count
charac	charac	bigint	bigint
P1	S1	900	2
P1	S5	100	1
P2	S5	300	2
P3	S2	3100	7
P3	S3	200	1
P3	S5	200	1
P4	S3	500	1
P4	S5	800	1
P5	S2	100	1
P5	S5	1000	3
P6	S4	600	2
P6	S5	700	2

Joseph Vella - (c) - SQL - Introduction

76

Operational view

Follow the steps for a possible operational view!

5 SELECT [DISTINCT] list of expressions

1 FROM list of tables

2 WHERE row level selection condition expression

3 [GROUP BY aggregate on expression list

4 [HAVING aggregate condition expression]]

ORDER BY list of attributes; 6

SELECT commission, COUNT (*)
FROM agents
GROUP BY commission
HAVING COUNT (*) > 3;

agents

AGENT_NAME	COMMISSION
Alex	.13
Subbarao	.14
Benjamin	.11
Ramasundar	.15
Alford	.12
Ravi Kumar	.15
Santakumar	.14
Lucida	.12
Anderson	.13
Mukesh	.11
McDien	.15
Ivan	.15

GROUP BY commission

COMMISSION	COUNT(*)
.15	4
.11	2
.14	2
.13	2
.12	2

HAVING COUNT (*) > 3;

COMMISSION	COUNT(*)
.15	4

COMMISSION	COUNT(*)
.11	2
.14	2
.13	2
.12	2

Joseph Vella - (c) - SQL - Introduction

77

Example

• See all "bins"!

```
SELECT wrk.qty,  
count(*)  
FROM date.spj wrk  
GROUP BY wrk.qty  
ORDER BY wrk.qty;
```

qty integer	count bigint
100	4
200	7
300	2
400	3
500	4
600	1
700	1
800	2

OK! But the boss wants bins from 0 to 1200!

```
SELECT bin.lp,  
count(*) "not goog!?",  
sum(CASE WHEN wrk.qty IS NULL THEN 0 ELSE 1 END) "that's it!"  
FROM date.spj wrk RIGHT OUTER JOIN  
generate_series(0, 1200, 100) as bin(lp)  
ON wrk.qty = bin.lp  
GROUP BY bin.lp  
ORDER BY bin.lp;
```

Outer join tuples have null!?

lp integer	not goog!? bigint	that's it! bigint
0	1	0
100	4	4
200	7	7
300	2	2
400	3	3
500	4	4
600	1	1
700	1	1
800	2	2
900	1	0
1000	1	0
1100	1	0
1200	1	0

Joseph Vella - (c) - SQL - Introduction

78

Another Example

- Which part is heavily used in works schedule (i.e. never used in quantity less than 400)?

```
SELECT wrk.p,  
       sum(wrk.qty) "Parts"  
FROM   date.spj wrk  
GROUP BY wrk.p  
HAVING min(wrk.qty) >= 400  
ORDER BY wrk.p;
```

p	Parts
character(20)	bigint
P4	1300

Joseph Vella - (c) - SQL - Introduction

79

Example

- Which job has the highest average product consumption?

```
SELECT wrk.j,  
       avg(wrk.qty)::integer  
FROM   date.spj wrk  
GROUP BY wrk.j  
HAVING avg(wrk.qty)  
       >= ALL (SELECT avg(wrk_in.qty)::integer  
               FROM date.spj wrk_in  
               GROUP BY wrk_in.j);
```

j	avg
character(20)	integer
J5	550

avg
integer
400
550
400
250
240
413
267

Joseph Vella - (c) - SQL - Introduction

80

Nested Queries

Includes sub-queries, correlated queries

Joseph Vella - (c) - SQL - Introduction

81

Introducing

- The nested genre allow us:
 - Defined universal and existential quantification queries:
 - **Existential type:**
 - Which parts are actually used?
 - **Universal type:**
 - Which jobs supplied by all suppliers?
- Nested queries tend to provide *data driven* capability to a query language.
- Nested queries in SQL are used extensively (i.e. not only in **SELECT** statements) – e.g. in **INSERT**, **UPDATE** ad **DELETE**.
 - In a **SELECT** statement a nested query plugs in a **WHERE** clause.
 - We have already seen some example with **IN** and **EXISTS** operators.

Joseph Vella - (c) - SQL - Introduction

82

The structure

General Structure of nested SELECT statement is:

Outer Query

```
SELECT [DISTINCT] list of attributes
FROM list of tables
WHERE [ row level selection condition expressions ] |
      [ row level comparison (
        SELECT [DISTINCT] list of attributes
        FROM list of tables
        WHERE row level selection condition expressions
      ) ] |
      [ row to aggregate level comparison (
        SELECT [DISTINCT] list of attributes
        FROM list of tables
        WHERE row level selection condition expressions
      ) ] ;
```

Inner Query

Notes:

- **Scoping Rules:**
 - Outer query **WILL NOT SEE** inner query rows!
 - An inner query row **will see**, unless name clash exists, the outer query's current row!
- **Latching** between outer and inner query needs our attention! (Use the right data type, number of attributes, and data type constructor (e.g. row IN (set of rows)).

Push Down tuple

Pop Result

83

Nested Query: Examples

Which part is not used in a work schedule?

```
SELECT prd.*
FROM date.p as p1d
WHERE prd.p NOT IN (SELECT DISTINCT p FROM date.spj);
```

	p	character(20)
1	P1	
2	P3	
3	P4	
4	P2	
5	P5	
6	P6	

	p	character(20)	pname	character(20)	colour	character(20)	weight	city	integer	character(20)
1	P7		LOCK NUT		GREY		13	PALO ALTO		

84

Nested Query: Examples

Which work schedule entry has the highest product usage?

```
SELECT wrk.*
FROM date.spj AS wrk
WHERE wrk.qty >= ALL
      (SELECT qty FROM date.spj);
```

s	character(20)	p	character(20)	j	character(20)	qty	integer
1	S2	P3	J7			800	
2	S5	P4	J4			800	

	qty	integer
1	200	
2	700	
3	400	
4	200	
5	200	
6	500	
7	600	
8	400	
9	800	
10	100	
11	200	
12	500	
13	300	
14	300	
15	200	
16	100	
17	500	
18	100	
19	200	
20	100	
21	200	
22	800	
23	400	
24	500	

Joseph Vella - (c) - SQL - Introduction

85

Nested Query: Examples

Which work schedule entry per job has the highest product usage?

```
SELECT wrk.*
FROM date.spj AS wrk
WHERE wrk.qty >= ALL
      (SELECT qty
       FROM date.spj
       WHERE j=wrk.j)
ORDER BY wrk.j;
```

Pushing 'j1'

Popping 200, 400

Final Result

s	character(20)	p	character(20)	j	character(20)	qty	integer
1	S2	P3	J1			400	
2	S3	P4	J2			500	
3	S4	P6	J3			300	
4	S5	P4	J4			800	
5	S2	P3	J5			600	
6	S2	P3	J6			400	
7	S2	P3	J7			800	

	qty	integer
	200	
	400	
	200	

This is an example of a co-related nested query!

Joseph Vella - (c) - SQL - Introduction

86

Which part is used in any work schedule?

```
SELECT prd.*
FROM date.p as prd
WHERE EXISTS (SELECT DISTINCT p
              FROM date.spj
              WHERE prd.p=p);
```

	p character(20)	pname character(20)	colour character(20)	weight integer	city character(20)
1	P1	NUT	RED	12	LONDON
2	P2	BOLT	GREEN	17	PARIS
3	P3	SCREW	BLUE	17	ROME
4	P4	SCREW	RED	14	LONDON
5	P5	CAM	BLUE	12	PARIS
6	P6	COG	RED	19	LONDON

Which part is not used in any work schedule (re-write)?

```
SELECT prd.*
FROM date.p as prd
WHERE NOT EXISTS (SELECT DISTINCT p
                  FROM date.spj
                  WHERE prd.p=p);
```

	p character(20)	pname character(20)	colour character(20)	weight integer	city character(20)
1	P7	LOCK NUT	GREY	13	PALO ALTO

Nested Query: Examples (UQ) I

- QUERY: "Which part is used in all jobs?"
- First, let us work out what "all jobs" is!
- Look for a data source that has the data (i.e. works!)
 - P, J, and WRK
- Examine the data
 - Clearly 'P3' works!
 - All other, e.g., 'P1', fail!?

```
SELECT *
FROM date.j AS job;
```

	p	j	name	city
	character (20)	character (20)	character (20)	character (20)
1	P1	J1	DISPLAY	ROME
2	P1	J2	OCR	ATHENS
3	P1	J3	CONSOLE	ATHENS
4	P1	J4	RAID	LONDON
5	P1	J5	EDS	OSLO
6	P1	J6	TAPE	LONDON

```
SELECT DISTINCT wrk.p, wrk.j
FROM date.spj AS wrk
ORDER BY wrk.p, wrk.j;
```

	p	j
	character (20)	character (20)
1	P1	J1
2	P1	J4
3	P2	J2
4	P2	J4
5	P3	J1
6	P3	J2
7	P3	J3
8	P3	J4
9	P3	J5
10	P3	J6
11	P3	J7
12	P4	J2
13	P4	J4
14	P5	J2
15	P5	J4
16	P5	J5
17	P5	J7
18	P6	J2
19	P6	J3
20	P6	J4
21	P6	J7

89

Nested Query: Examples (UQ) II

- QUERY (continued): "Which part is used in all jobs?"
- How to work out success and fail query for 'P1' & 'P3'?

```
SELECT *
FROM DATE.J AS JOB
WHERE NOT EXISTS
  (SELECT *
   FROM DATE.SPJ AS WRK
   WHERE WRK.P = 'P1'
   AND WRK.J = JOB.J);
```

	j	jname	city
	[PK] character (20)	character (20)	character (20)
1	J2	DISPLAY	ROME
2	J3	OCR	ATHENS
3	J5	RAID	LONDON
4	J6	EDS	OSLO
5	J7	TAPE	LONDON

```
SELECT *
FROM DATE.J AS JOB
WHERE NOT EXISTS
  (SELECT *
   FROM DATE.SPJ AS WRK
   WHERE WRK.P = 'P3'
   AND WRK.J = JOB.J);
```

Data OutputMessagesNotifications

	j	jname	city
	[PK] character (20)	character (20)	character (20)

No jobs - OK!

Total rows: 0 of 0 Query complete 00:00:00.057

90

Nested Query: Examples
(UQ) III

Which part is used in all jobs?

```
SELECT DISTINCT prd.p
FROM date.p AS prd
WHERE NOT EXISTS ( SELECT *
                   FROM date.j AS job
                   WHERE NOT EXISTS ( SELECT *
                                     FROM date.spj AS wrk
                                     WHERE wrk.p=prd.p
                                       AND wrk.j=job.j));
```

	p
	character(20)
1	P3

This is an example of a **Universal Quantification Query (UQ)**!

Joseph Vella - (c) - SQL - Introduction

91

Views

Are a convenient way to organise and rehash database data.

Joseph Vella - (c) - SQL - Introduction

92

Data Views

- Data views are aspects, perspective of a database portion.
 - We specify the structure and possible content of view through a query (e.g. a SELECT statement).
 - Content is generated on demand or pre-computed.
- Views are ideal for retrieving data.
 - Reports;
 - Forms;
 - Generating structure and content of datasets.
- Views have come a long way and these are a basis for:
 - Parameterised views;
 - Remote views;
 - Materialised views;
 - Updateable views.
- The following is the generic syntax; note that query could be any query we have developed here.

```
CREATE [ OR REPLACE ] [ TEMP | TEMPORARY ] [ RECURSIVE ] VIEW name [ ( column_name [, ...] ) ]
[ WITH ( view_option_name [= view_option_value] [, ...] ) ]
AS query
[ WITH [ CASCADED | LOCAL ] CHECK OPTION ]
```

Joseph Vella - (c) - SQL - Introduction

93

Issues with views!?

- But can we use them to manipulate the underlying data?
 - In general no.
 - What to do when a view's query is based on a:
 - JOINS are involved (even lossy ones);
 - GROUP BY queries;
 - Use of DISTINCT keyword;
 - LIMIT & OFFSET keywords;
 - UNION etc;
 - Missing attributes (e.g. primary key attribute and other constraints);
 - Etc
 - THIS IS CALLED THE VIEW UPDATE PROBLEM!
 - Recent advances in SQL actually allow to address some of the above issues:
 - Based on attaching code to a view and this is invoked when one tries to manipulate data through the view.

Joseph Vella - (c) - SQL - Introduction

94

Management of a view

- Define it!
 - CREATE VIEW

```
CREATE OR REPLACE VIEW date.top_s AS
SELECT sup.*
FROM date.s AS sup
WHERE status = 10;

CREATE OR REPLACE VIEW
date.top_s(topsupplier, hiscity) AS
SELECT sname, city
FROM date.s AS sup
WHERE status = 10;
```

- Use it. (And re-use it!)

```
SELECT *
FROM date.top_s;
```

	topsupplier character(20)	hiscity character(20)
1	JONES	PARIS

```
SELECT wrk.p, wrk.j
FROM date.spj AS wrk INNER JOIN
date.top_s AS top ON wrk.s = top.s;
```

- Drop it!?

```
DROP VIEW date.top_s;
```

Joseph Vella - (c) - SQL - Introduction

95

Views are visible in the DD

The screenshot shows the pgAdmin III interface. On the left, the 'Object browser' displays the database structure, with 'Views (1)' highlighted under the 'date' schema. The 'Properties' pane on the right shows details for the view 'date.top_s', including its definition: 'SELECT sup.sname AS topsupplier, sup.city AS hiscity FROM date.s sup WHERE status = 10;'. The 'SQL pane' at the bottom contains SQL commands to create, drop, and alter the view.

Joseph Vella - (c) - SQL - Introduction

96

User Defined Functions (UDFs)

We are allowed to write our own functions!

The programming language could one of a few, e.g. Java, PL/pgsql, and SQL.

Joseph Vella - (c) - SQL - Introduction

97

Basics

- User defined functions allows the database user to extend the functionality of SQL over his data.
- Functions have a name (e.g. can include a schema name too).
 - A function can take zero to many named and typed arguments;
 - A function returns an object:
 - Row;
 - Void (nothing!?!);
 - Set of rows (e.g. a table).
 - A function's body is a sequence, if language is SQL, of SQL statements.
 - The last statement must return values if its return type is not void.
- If the function writing language is SQL then the result of a function is based on the last executed.
 - If the result type of the function is a row, the *first* row computed for the last query is returned.
 - Beware! This is not well defined ...
- An advanced note:
Up to early versions of PostgreSQL 10, it is not possible to commit new transactions in a function – i.e. no autonomous transaction mechanism exist.
 - To work around this, the function, to commit an autonomous transaction, must be run in a new session – invoke a function on another server connection!

Joseph Vella - (c) - SQL - Introduction

98

Define and use a UDF

Given a job, as an argument, work out the relative entries in the work schedule.

```
CREATE OR REPLACE
    FUNCTION date.getSupByJob(jpk text)
    RETURNS integer AS $$
        SELECT count(*)::integer
        FROM date.spj
        WHERE j=jpk;
    $$ LANGUAGE SQL;
```

← Execute this to compile it!

```
-- usage examples

SELECT date.getSupByJob('J1');
-- 3::integer

SELECT job.*
FROM date.j AS job
WHERE date.getSupByJob(job.j)>=3;
```

	j	character(20)	jname	character(20)	city	character(20)
1	J1		SORTER		PARIS	
2	J2		DISPLAY		ROME	
3	J4		CONSOLE		ATHENS	
4	J7		TAPE		LONDON	

Joseph Vella - (c) - SQL - Introduction

99

An example

Get all work detail records related to job 'J1'.

```
CREATE OR REPLACE
    FUNCTION date.getWorksSupByJob(jpk text)
    RETURNS setof date.spj AS $$
        SELECT *
        FROM date.spj
        WHERE j=jpk;
    $$ LANGUAGE SQL;
```

← Execute this to compile it!

```
SELECT date.getWorksSupByJob('J1');

SELECT *
FROM date.getWorksSupByJob('J1');
```

	s	character(20)	p	character(20)	j	character(20)	qty	integer
1	S1		P1		J1		200	
2	S2		P3		J1		400	
3	S3		P3		J1		200	

Joseph Vella - (c) - SQL - Introduction

100

- How to implement user defined views (e.g. views with a parameter)?
 - First create a view (with no data restriction);
 - Create a function to access the view and apply a restriction.
 - Execute the function with particular value to restrict on.

```
-- sort of dynamic view on the following
SELECT p, count(*),min(qty),max(qty)
FROM date.spj
GROUP BY p;
```

	p	count	min	max
	character(20)	bigint	integer	integer
1	P1	3	100	700
2	P3	9	200	800
3	P4	2	500	800
4	P2	2	100	200
5	P5	4	100	500
6	P6	4	200	500

```
CREATE OR REPLACE VIEW date.p_stat_spj AS
SELECT p, count(*),min(qty),max(qty)
FROM date.spj
GROUP BY p;
```

```
SELECT * FROM date.p_stat_spj_p('P1');
```

```
CREATE OR REPLACE
FUNCTION date.p_stat_spj_p(ppk text)
RETURNS date.p_stat_spj AS $$
SELECT *
FROM date.p_stat_spj
WHERE p=ppk;
$$ LANGUAGE SQL;
```

	p	count	min	max
	character(20)	bigint	integer	integer
1	P1	3	100	700