



# Introduction to Databases

**Prof. Joseph G. Vella**  
**Dept. of Computer Information Systems**  
JOSEPH.G.VELLA@UM.EDU.MT


1



# Concurrency & Transaction Processing

**Rationale & Challenges**

2




## Where had we started ...

- **Definition:**
  - A **database** is a collection of **structured data** stored on **persistent storage**.
- **Environment and Contextual**
  - **Available** to a large number of end users.
    - As in “Access my data when I need it ...”
      - **No delay, no access denial, no loss of data**
  - **Correctness**.
    - As in “What’s my balance?” gets the same value until it is changed.
- **What’s different from, say, programming?**
  - Data **persistence**;
  - Data **sharing**;
  - Data **independence**.

Joseph G Vella - Introduction to databases 3 Concurrency & Recovery

3




## How is usage managed and controlled?

- **A Database Management System (DBMS)**
  - Is a computer-based application tool set for:  
**defining, creating, manipulating, controlling, and managing** databases.
- **DBMS Requirements:**  
End users, developers, administrators demand **efficiency** in terms of allocation of computational resources and storage space.
  - **Reliability**
  - **Openness** (e.g., in term of data connectivity)
  - **Scalability** (e.g., storage and throughput)

Joseph G Vella - Introduction to databases 4 Concurrency & Recovery

4




## DBMS Functionality

- DBMS also facilitate and insulates data access by system's end users (including "computer" naive *end users*).
- Three important functionalities include:
  - *Query Processing*
  - *Transaction Management*
  - *Storage management*
- Two main areas of concern:
  - *Consistency* (e.g. returns the same results under invariant states), and
  - *Efficiency* (e.g. computational time and space) of DBMS activities.
- Some advantages:
  - Execute a query in an efficient and opportunistic approach – *query optimisation*;
  - Iron out sharing inconsistencies between conflicting queries and transactions – *concurrency control*;
  - Undoing the effects of incomplete transactions under a system's instructions or in rectifying a system failure – *recovery*.
    - Note: system here is most general: e.g., OS, DBMS, hard disk

Joseph G Vella - Introduction to databases 5 Concurrency & Recovery

5



## Transaction Processing

6




## Transactions & Transaction Processing

- A **transaction** is a sequence of actions over a database that realise a logical operation.
  - For example, a double entry accounting package must execute the debit and the credit entries for the database to remain consistent;
    - In this case the transaction is made up of two SQL insert statements.
- The part of a DBMS that applies transaction to its database state is called the **Transaction Processing** module.
- Basic actions over a database are:
  - **read(X)** -- read item X from database and store into local var X; and
  - **write(X)** -- write item X to database from local var X.
- Concurrency control is that part of the DBMS that **assures** that simultaneously executed transactions produce the same results as if they were executed in serially.
- The concurrency control mechanism must constantly seek to **balance** a high degree of concurrency against a satisfactory performance within a consistent state.



## TP Principles - ACID

- Transaction's **ACID** principles for **short-lived** transactions are (over and above valid transactions):
  - **Atomicity**
    - The transaction executes completely or not at all,
  - **Consistency**
    - The transaction preserves the internal consistency of the database.
  - **Isolation**
    - The transaction executes as if it were running alone, with no other transactions.
  - **Durability**
    - The transaction's result will not be lost in the future.




## Transaction Modeling

- An explicit sequence of data manipulation commands over a database are abstracted as a single logical update (i.e a **transaction**).
  - A transaction is **valid** if **all** its sub-parts (e.g. updates, deletes) have succeeded. A valid transaction can **commit**, otherwise we have to **rollback** the transaction.
  - It is possible to have **read-only transactions** apart from read-write transactions.
- A number of basic TP management operations are required (to supplement READ & WRITE basic ops):
  - BEGIN, END, Request a COMMIT, ROLLBACK;
  - UNDO and REDO specific log entries (e.g. READ & WRITE based).
- Operationally, this complicates our scenario when allowing for data sharing! The DBMS typically insulates this from users (programmers and end users alike!).

Joseph G Vella - Introduction to databases
9
Concurrency & Recovery

9




## A Database Transaction : an Example

- “Ensure that the employee's salary is upped by 25% and each salesperson commission is augmented by Euro 250.”
- A generic SQL command sequence :
 

```
START transaction;
UPDATE emp
    SET sal = sal * 1.25;
Remark IMPLICIT mechanism - ON error ROLLBACK
UPDATE emp
    SET comm = comm + 250 WHERE comm IS NOT NULL;
Remark IMPLICIT mechanism - ON error ROLLBACK
COMMIT;
END transaction;
```

Joseph G Vella - Introduction to databases
10
Concurrency & Recovery

10



### TP & a sequence of SQL DML statements

- Here is a sequence of generic SQL statements that are entered consecutively. We will assume that a `START TRANSACTION` statement at the beginning of each transaction is implicit.


%	SQL Construct	Explanation
1	INSERT ...	Ineffective as regards persistence.
2	DELETE ...	<i>ditto</i>
3	SAVEPOINT sp1	Define savepoint sp1.
4	INSERT ...	Ineffective as regards persistence.
5	SAVEPOINT sp2	Define savepoint sp2.
6	DELETE ...	Ineffective as regards persistence.
7	ROLLBACK WORK TO SAVEPOINT SP2	All DML since <i>sp2</i> are un-done (i.e. 6).
8	UPDATE ...	Ineffective as regards persistence.
9	ROLLBACK WORK TO SAVEPOINT SP1	All DML since <i>sp1</i> are un-done (i.e. 4 & 8)
10	DELETE ...	Ineffective as regards persistence.
11	UPDATE ...	Ineffective as regards persistence.
12	COMMIT WORK	All DML since the start become persistent (i.e. 1,2,10 & 11).
13	INSERT ...	Ineffective as regards persistence.
14	quit	All pending DML since the start are un-done (i.e. 13).

Joseph G Vella - Introduction to databases

11

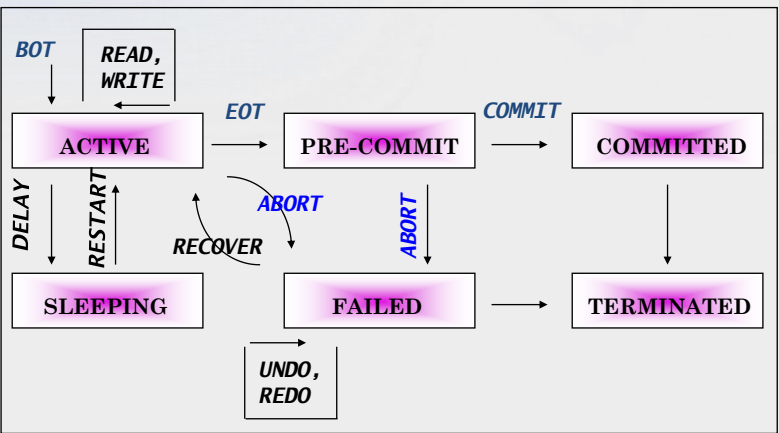
Concurrency & Recovery

11



### TP State Transition Model

A state transition diagram for a transaction (and recovery) sub-system.



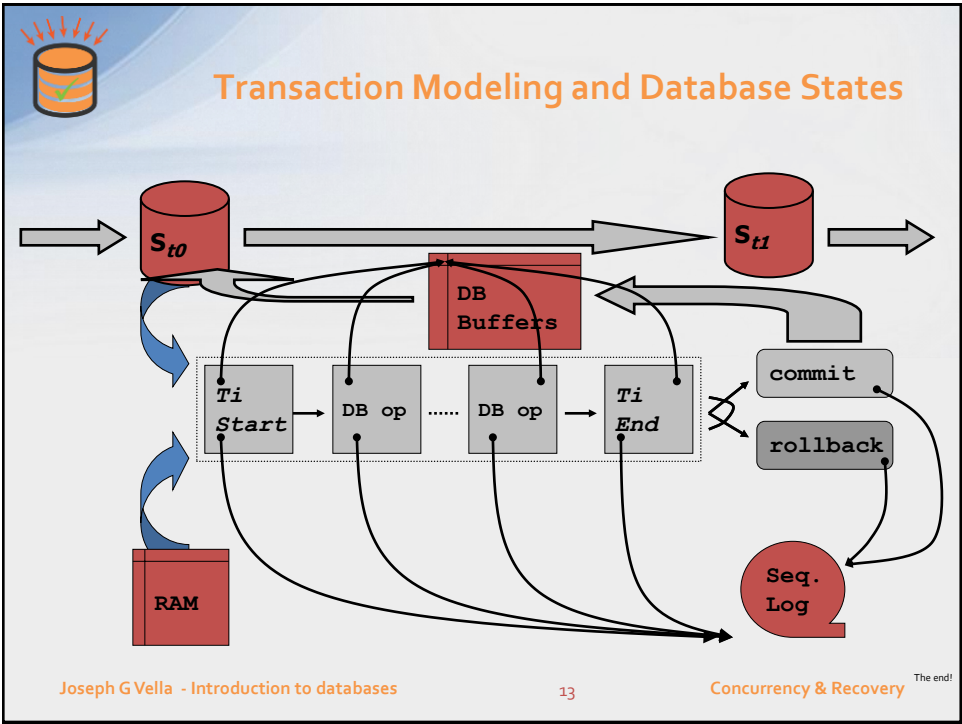
```
graph LR
    BOT[ ] --> ACTIVE
    subgraph Box1 [ ]
        ACTIVE[ACTIVE]
        SLEEPING[SLEEPING]
        FAILED[FAILED]
    end
    subgraph Box2 [ ]
        PRECOMMIT[PRE-COMMIT]
        COMMITTED[COMMITTED]
        TERMINATED[TERMINATED]
    end
    ACTIVE -- EOT --> PRECOMMIT
    PRECOMMIT -- COMMIT --> COMMITTED
    ACTIVE -- DELAY --> SLEEPING
    SLEEPING -- RESTART --> ACTIVE
    ACTIVE -- ABORT --> FAILED
    FAILED -- RECOVER --> PRECOMMIT
    FAILED -- ABORT --> TERMINATED
    COMMITTED --> TERMINATED
    FAILED -- UNDO, REDO --> UNDO[UNDO, REDO]
```

Joseph G Vella - Introduction to databases

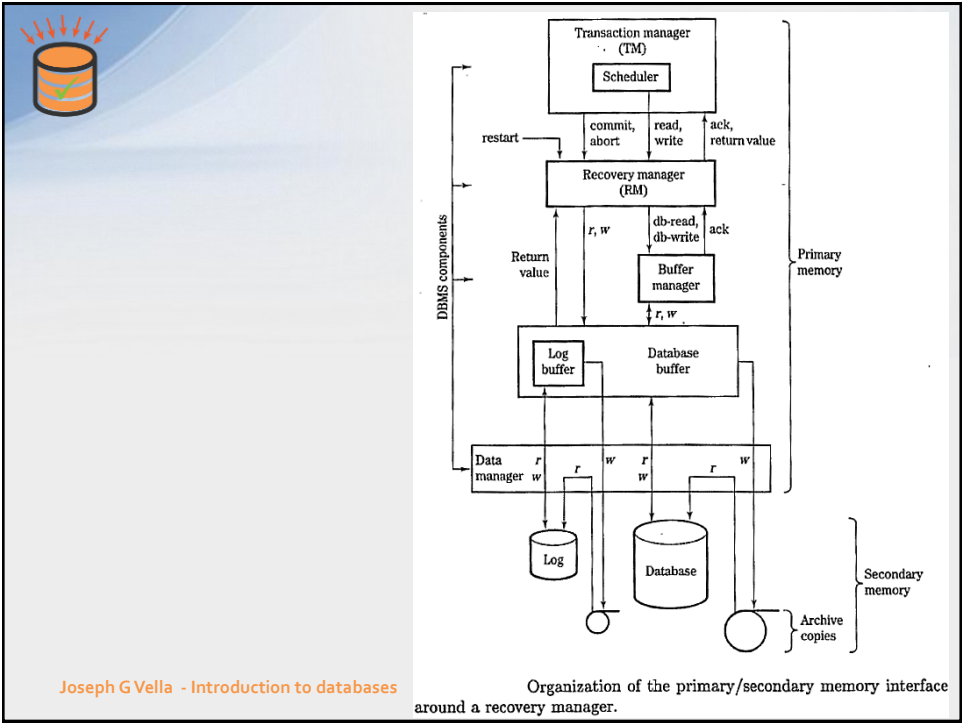
12

Concurrency & Recovery


12



13



14



## SQL Transaction Management


- SQL provides constructs for the management of transactions:
  - Starting a transaction
    - `START TRANSACTION` #in some dialects!
  - Set a transaction sharing level (per session or transaction)
    - `SET TRANSACTION ISOLATION LEVEL`  
{ `READ {UNCOMMITTED | COMMITTED } |`  
`REPEATABLE READ |`  
`SERIALIZABLE |`  
`SERIALIZABLE }`;
  - Committing a transaction (i.e. Requesting a Commit)
    - `COMMIT`
      - To make permanent the changes made by statements issued and the beginning of a transaction.
  - Save points
    - `SAVEPOINT sp4`
      - To establish a point back to which you may roll.
  - Rolling back transactions
    - `ROLLBACK TO`  
`SAVEPOINT sp4`
      - To undo all changes since the beginning of a transaction or some savepoint.

Joseph G Vella - Introduction to databases

15

Concurrency & Recovery

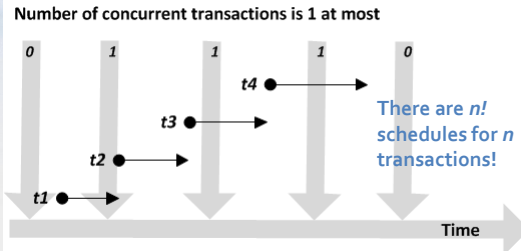
15



## Serial vs. Interleaving Execution

Serial

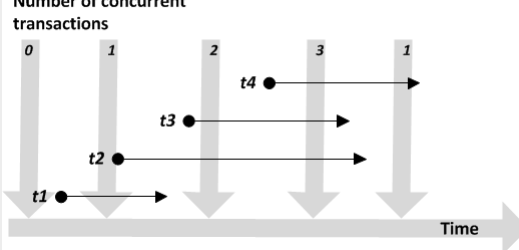
Number of concurrent transactions is 1 at most



There are  $n!$  schedules for  $n$  transactions!

Interleaving

Number of concurrent transactions



There are  $(\sum(tn_j))/n!$  schedules for  $n$  transactions (with  $tn_j$  representing the number of R/W ops of transaction  $tn$ ! i.e., much much more!?


Joseph G Vella - Introduction to databases

16

Concurrency & Recovery

16





### Unrestricted access to a common database state: *dirty read*

- An unrestricted sharing over data is problematic:


Time	Database state	Process & Trans. 1		Process and Trans. 2		Comments
		Transaction	Local memory	Transaction	Local memory	
			X=0		Y=0	
1	A=10, & B=20	read(A,X)	X=10			Read db rec A to var X.
2	A=10, & B=20	X:=X+100	X=110			
3	A=10, & B=20	write(A,X)	X=110			Write var X to db A.
4	A=110, & B=20			read(A,Y)	Y=110	T2 read from T1, and T1 is not yet closed.
5	A=110, & B=20			write(B,Y)	Y=110	
6	A=110, & B=110	rollback	x=10			
7	A=10, & B=110					At this point B is incorrect

Joseph G Vella - Introduction to databases

17

Concurrency & Recovery

17



### Unrestricted access to a common database state: *lost update*

- Concurrency problem - Lost operations
  - This problem arises due to the following sequence of events;


Time	Tran. T1	Tran. T2
t1	Read(X)	
t2		Read(X)
t3		X:=(X+1)
t4		Write(X)
t5	X:=(X+1)	
t6	Write(X)	

Joseph G Vella - Introduction to databases

18

Concurrency & Recovery

18



### Unrestricted access to a common database state: *inconsistent read*

- **Concurrency problem -- Inconsistencies**
  - This is due to the presence of an integrity constraint. The problem arises each time a transaction accesses or modifies a transitional state of the database characterised by the fact that an integrity constraint is not verified.
  - I.C.  $A = B$ ;


<u>Time</u>	<u>Tran. T1</u>	<u>Tran. T2</u>	
t1	Read(A)		
t2	A:=A+1		
t3	Write(A)		
t4		Read(A)	
t5		Print(A)	<i>ic broken <math>A \neq B</math></i>
t6		Read(B)	
t7		Print(B)	<i>ic broken <math>A \neq B</math></i>
t8	Read(B)		
t9	B:=B+1		
t10	Write(B)		

Joseph G Vella - Introduction to databases

19

Concurrency & Recovery

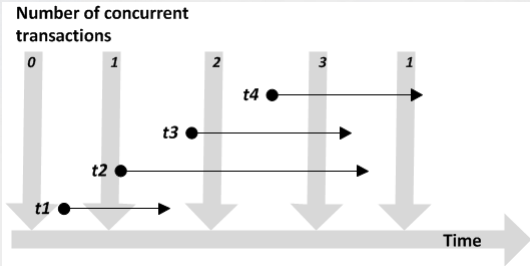
19

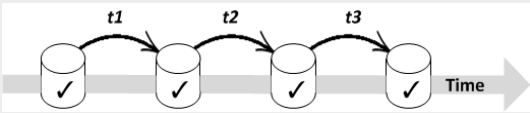


### Sharing and Consistency: Aims

- **Ideally the DBMS must support:**
  - High level of sharing
    - That is, the more concurrent users are accessing the database the better.
  - Consistency enforcement
    - But only at the end of each transaction consistency is maintained.

Number of concurrent transactions






Joseph G Vella - Introduction to databases

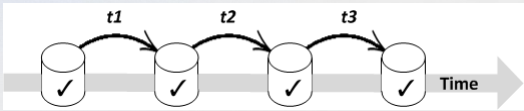
20

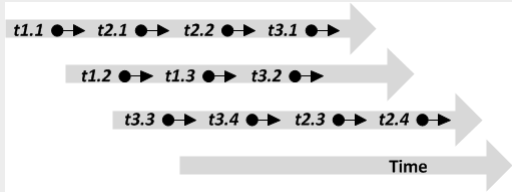
Concurrency & Recovery

20



Is (there any) Serial and its Interleaved transaction resulting in the same database end state?





If yes, then that interleaving as the same effect as the serial!

Joseph G Vella - Introduction to databases

21

Concurrency & Recovery

21



TP & Concurrent Control:  
Two Phase Locking

22



## Two Phase Locking (2PL)

- Is there an algorithm that can ensure concurrent transactions over a database have the same affect as a (defined) sequence of serial transactions?
  - Remember the extreme points:
    - Serial means correct but lowest sharing;
    - Interleaved means better utilization through higher access but possible sharing violations (if unrestricted).
- Yes for centralised data servers and “short” transactions we use the *two phase locking* algorithm.


23



## Locking

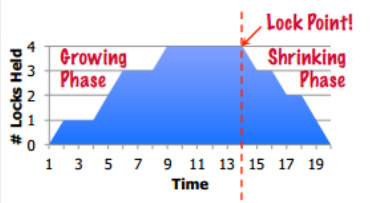
- Locks are variables linked with instances and holds the state of the instances acceptable access operations. Most DBMS have a *lock manager*.
  - Types of locks include:
    - *binary locks*
      - an item can either be LOCKED or UNLOCKED.
        - ideal for *mutual exclusion*.
      - typical operations of a binary lock are:
        - lock(X);
        - unlock(X).
      - » THESE OPERATIONS MUST BE INDIVISIBLE.
    - for the transaction model a lock operation is invoked when:
      - lock(X) before every read(X) or write(X) of T<sub>i</sub>;
      - unlock(X) after all r/w s are completed in T<sub>i</sub>;
      - no issues of lock(X) if it already holds a lock on X;
      - no issues of unlock(X) if it already holds X.
- Locking on its own does not solve concurrent access violation – unless we don't apply serial transaction processing (i.e. lock the whole database for each transaction in turn).
- Heuristic: *where there are locks there is deadlocks!?*

24



2PL

- It is the 2pl scheduler job to manage the locks by controlling when transactions get and free locks.
- 2pl protocols are characterised by dividing locking operations into two phases: the first being the locks hoarding (*growing*) phase followed by a *shrinking* (releasing) of locks phase.
  - During the growing phase *no locks are released*.
  - During the shrinking phase *no locks are requested*.




- Example:
  - T1: r1k(Y), r(Y), u1k(Y), w1k(X), r(X), ..., w(X), u1k(X).
  - T1's schedule does not follow the 2pl protocol because
    - ulk(Y) precedes wlk(X)!
  - T2: r1k(Y), r(Y), w1k(X), u1k(Y), r(X), ..., w(X), u1k(X).
  - T2's schedule does follow a 2pl protocol!

Joseph G Vella - Introduction to databases

25

Concurrency & Recovery

25



2PL protocols can help with serializability


- There is an *important* theorem that states:
  - If all transactions in a schedule adhere to the 2PL protocol then the schedule is serialisable.
    - no need to test for serialisability of schedules;
    - the protocol enforces serialisability;
    - but limits the level of concurrency in a schedule
      - Why?

Joseph G Vella - Introduction to databases

26

Concurrency & Recovery

26




## 2PL variants

- Other types of 2pl (the previous one is called basic 2pl):
  - *conservative 2pl*
    - Hoard all locks at the beginning of a transaction.
      - If any *lock* (on a read or write of the transaction) cannot be accessed it waits until all required locks are available.
    - This protocol is of course **deadlock-free!**
  - *strict 2pl*
    - Release all of a transaction's locks on commit.
    - This protocol is *not* deadlock-free!?
- The scheduler needs a strategy for *detecting deadlocks* (i.e. no transaction is indefinitely blocked).
  - A simple mechanism is by a **time-out span**. If a transaction has been waiting *more* than the set time-out for a lock it assumes there is a deadlock and aborts (usually, the transaction longest in waiting).


Joseph G Vella - Introduction to databases      27      Concurrency & Recovery

27



## SQL Isolation vs Data Sharing Issues

28



### SQL Isolation vs Data Sharing Issues

Set a transaction sharing level (per session or transaction)

```
SET TRANSACTION ISOLATION  
LEVEL  
{ READ {UNCOMMITTED |  
COMMITTED } |  
REPEATABLE READ |  
SNAPSHOT |  
SERIALIZABLE };
```

**Very Important:**  
This is how the SQL  
Standard see  
management of TP;  
Each vendor adopts with  
its own mechanism. This  
leads to some  
interpretation issues.

Isolation Level	Phenomena Allowed
Read uncommitted	Dirty Read Non Repeatable Read Phantoms
Read committed	Non Repeatable Read Phantoms
Repeatable read	Phantoms
Serializable	-

Concurrence ↑ Consistency ↓

Joseph G Vella - Introduction to databases

29

Concurrency & Recovery

29






Data  
Recovery

Database Recovery

30



### Database Recovery


- DBMSs must ensure that a transaction either succeeds or else, on failure, the database is kept in a consistent state (i.e. this process is called recovery).
- What are the types of failures the DBMSs' transaction sub-system must recover from:
  - computer hardware, software or network related failures (external to the DBMSs activities)
    - can affect main and secondary storage structures held by the DBMS;
  - the DBMSs sub-system that handles transaction processing can have an operational failure
    - divide by zero during the processing of a transaction;
    - illogical sequence of coding in the transaction or system;
  - the transaction cannot continue due to conflicts with the present data
    - cannot deduct money from an account that is closed!;
  - concurrency enforcement problems;
  - disk, network, etc, failure during a DBMS activity;
  - accidental / intentional factors (earthquakes, viruses, deletion of database entities by the system administrator).

Joseph G Vella - Introduction to databases

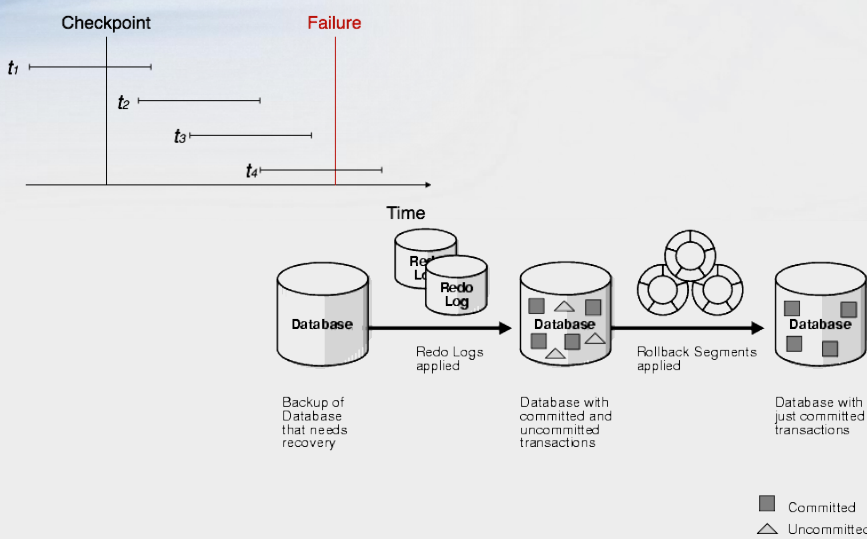
31

Concurrency & Recovery

31



### We need to restore database state as at end of $t_3$



The diagram illustrates the recovery process. A timeline at the top shows transactions  $t_1, t_2, t_3, t_4$  and a failure occurring at the end of  $t_3$ . Below, the recovery steps are shown: 1. Backup of Database that needs recovery. 2. Redo Logs applied, resulting in a database with committed (squares) and uncommitted (triangles) transactions. 3. Rollback Segments applied, resulting in a database with just committed transactions. A legend indicates that squares represent committed transactions and triangles represent uncommitted transactions.

Joseph G Vella - Introduction to databases

32

Concurrency & Recovery

32





## Recovery Process

- **Recovery entails:**
  - *detection* of a problem; and
  - *co-ordinated action* to eliminate the problem and its effects.
- **How to recover?**
  - Clearly the errors just mentioned require different *strategies* and *tactics*.  
But the worst are those that are associated with a *loss of data*.
  - Most recovery algorithms have these two components:
    - precautions taken during normal transaction processing to provide for a high level of state recovery;
    - actions taken after a failure to ensure the database recovery is graceful and effective.
  - By building a matrix of storage types vs failure types one can distinctly visualise the variety and extent of recovery procedures: